

More SQL: Complex Queries, Views, and Schema Modification

More Complex SQL Retrieval Queries

- Additional features allow users to specify more complex retrievals from database:
 - Nested queries, joined tables, and outer joins (in the FROM clause), aggregate functions, and grouping

Schema Change Statements in SQL

- **Schema evolution commands**

- DBA may want to change the schema while the database is operational
- Does not require recompilation of the database schema

The DROP Command

- DROP command
 - Used to drop named schema elements, such as tables, domains, or constraint
- Drop behavior options:
 - CASCADE and RESTRICT
- Example:
 - `DROP SCHEMA COMPANY CASCADE;`
 - This removes the schema and all its elements including tables, views, constraints, etc.

The ALTER table command

- **Alter table actions** include:
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints
- **Example:**
 - `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job
VARCHAR(12);`

Alter Statement

Alter Statement:

used to make changes to the schema of the table. Columns can be added and the data type of the columns changed as long as the data in those columns conforms to the data type specified.

Syntax:

```
ALTER TABLE table_name  
ADD (column datatype [Default Expression])  
[REFERENCES table_name (column_name)]  
[CHECK condition]
```

- **Example:**

```
ALTER TABLE studios  
ADD (revenue Number int )
```

Alter Statement

Add table level constraints:

Syntax:

```
ALTER TABLE table_name  
ADD ([CONSTRAINT constraint_name CHECK comparison]  
[columns REFERENCES table_name (columns)])
```

Example:

```
ALTER TABLE studios  
ADD (CONSTRAINT check_state CHECK (studio_state in ('TX', 'CA', 'WA')))
```

Modify Columns:

Syntax:

```
ALTER TABLE table_name  
MODIFY column [data type]  
[Default Expression]  
[REFERENCES table_name (column_name)]  
[CHECK condition]
```

Example:

```
ALTER TABLE People  
MODIFY person_union varchar(10)
```

- **Notes1:** Columns can not be removed from the table using alter. If you want to remove columns you have to drop the table and then recreate it without the column that you want to discard

Adding and Dropping Constraints

- Change constraints specified on a table
 - Add or drop a named constraint

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```


Dropping Columns, Default Values

- To drop a column
 - Choose either CASCADE or RESTRICT
 - CASCADE would drop the column from views etc. RESTRICT is possible if no views refer to it.

ALTER TABLE COMPANY.EMPLOYEE DROP COLUMN Address CASCADE;

- Default values can be dropped and altered :

ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn DROP DEFAULT;

ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn SET DEFAULT '333445555';

Specifying Joined Tables in the FROM Clause of SQL

- **Joined table**

- Permits users to specify a table resulting from a join operation in the FROM clause of a query
- The FROM clause in Q1A
 - Contains a single joined table. JOIN may also be called INNER JOIN

```
Q1A:  SELECT  Fname, Lname, Address
      FROM    (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE   Dname='Research';
```

Different Types of JOINed Tables in SQL

- Specify different types of join
 - NATURAL JOIN
 - Various types of OUTER JOIN (LEFT, RIGHT, FULL)
- NATURAL JOIN on two relations R and S
 - No join condition specified
 - Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

NATURAL JOIN

- Rename attributes of one relation so it can be joined with another using NATURAL JOIN:

```
Q1B:   SELECT    Fname, Lname, Address
FROM      (EMPLOYEE NATURAL JOIN
              (DEPARTMENT AS DEPT (Dname, Dno, Mssn,
                                     Msdate)))
WHERE     Dname='Research';
```

The above works with $EMPLOYEE.Dno = DEPT.Dno$ as an implicit join condition

INNER and OUTER Joins

- INNER JOIN (**versus** OUTER JOIN)
 - Default type of join in a joined table
 - Tuple is included in the result only if a matching tuple exists in the other relation
- LEFT OUTER JOIN
 - Every tuple in left table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of right table
- RIGHT OUTER JOIN
 - Every tuple in right table must appear in result
 - If no matching tuple
 - Padded with NULL values for attributes of left table

Example: LEFT OUTER JOIN

```
SELECT E.Lname AS Employee_Name  
       S.Lname AS Supervisor_Name
```

```
FROM Employee AS E LEFT OUTER JOIN EMPLOYEE AS S  
      ON E.Super_ssn = S.Ssn)
```

ALTERNATE SYNTAX:

```
SELECT E.Lname , S.Lname  
FROM EMPLOYEE E, EMPLOYEE S  
WHERE E.Super_ssn + = S.Ssn
```

Multiway JOIN in the FROM clause

- FULL OUTER JOIN – combines result if LEFT and RIGHT OUTER JOIN
- Can nest JOIN specifications for a multiway join:

```
Q2A: SELECT Pnumber, Dnum, Lname, Address, Bdate  
          FROM ((PROJECT JOIN DEPARTMENT ON  
              Dnum=Dnumber) JOIN EMPLOYEE ON  
Mgr_ssn=Ssn)  
          WHERE Plocation='Stafford';
```

Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
 - **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- **Grouping**
 - Create subgroups of tuples before summarizing
- To select entire groups, **HAVING** clause is used
- Aggregate functions can be used in the **SELECT** clause or in a **HAVING** clause

Renaming Results of Aggregation

- Following query returns a single row of computed values from EMPLOYEE table:

**Q19: SELECT SUM (Salary), MAX (Salary), MIN (Salary), AVG
 (Salary)
 FROM EMPLOYEE;**

- The result can be presented with new names:

```
Q19A:      SELECT SUM (Salary) AS Total_Sal, MAX (Salary) AS
            Highest_Sal, MIN (Salary) AS Lowest_Sal, AVG
            (Salary)
            AS Average_Sal
            FROM    EMPLOYEE;
```

Aggregate Functions in SQL (cont'd.)

- NULL values are discarded when aggregate functions are applied to a particular column

Query 20. Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
Q20:  SELECT    SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
      FROM      (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)
      WHERE     Dname='Research';
```

Queries 21 and 22. Retrieve the total number of employees in the company (Q21) and the number of employees in the 'Research' department (Q22).

```
Q21:  SELECT    COUNT (*)
      FROM      EMPLOYEE;
```

```
Q22:  SELECT    COUNT (*)
      FROM      EMPLOYEE, DEPARTMENT
      WHERE     DNO=DNUMBER AND DNAME='Research';
```

Aggregate Functions on Booleans

- SOME and ALL may be applied as functions on Boolean Values.
- SOME returns true if at least one element in the collection is TRUE (similar to OR)
- ALL returns true if all of the elements in the collection are TRUE (similar to AND)

Grouping: The GROUP BY Clause

- **Partition** relation into subsets of tuples
 - Based on **grouping attribute(s)**
 - Apply function to each such group independently
- **GROUP BY** clause
 - Specifies grouping attributes
- **COUNT (*)** counts the number of rows in the group

Examples of GROUP BY

- The grouping attribute must appear in the SELECT clause:

```
Q24:          SELECT      Dno, COUNT (*), AVG (Salary)
              FROM        EMPLOYEE
              GROUP BY    Dno;
```

- If the grouping attribute has NULL as a possible value, then a separate group is created for the null value (e.g., null Dno in the above query)
- GROUP BY may be applied to the result of a JOIN:

```
Q25:  SELECT      Pnumber, Pname, COUNT (*)
      FROM        PROJECT, WORKS_ON
      WHERE              Pnumber=Pno
      GROUP BY    Pnumber, Pname;
```

Grouping: The GROUP BY and HAVING Clauses (cont'd.)

- **HAVING** clause
 - Provides a condition to select or reject an entire group:
- **Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

```
Q26:          SELECT      Pnumber, Pname, COUNT (*)
               FROM        PROJECT, WORKS_ON
               WHERE        Pnumber=Pno
               GROUP BY     Pnumber, Pname
               HAVING       COUNT (*) > 2;
```

Combining the WHERE and the HAVING Clause

- Consider the query: we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.

- INCORRECT QUERY:**

```
SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary>40000
GROUP BY    Dno
HAVING      COUNT (*) > 5;
```

Combining the WHERE and the HAVING Clause (continued)

Correct Specification of the Query:

- Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

Query 28. For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

```
Q28:  SELECT  Dnumber, COUNT (*)
      FROM    DEPARTMENT, EMPLOYEE
      WHERE   Dnumber=Dno AND Salary>40000 AND
            ( SELECT      Dno
              FROM        EMPLOYEE
              GROUP BY Dno
              HAVING      COUNT (*) > 5)
```


Views (Virtual Tables) in SQL

- Concept of a view in SQL
 - Single table derived from other tables called the **defining tables**
 - Considered to be a virtual table that is not necessarily populated

Specification of Views in SQL

- **CREATE VIEW** command

- Give table name, list of attribute names, and a query to specify the contents of the view
- In V1, attributes retain the names from base tables. In V2, attributes are assigned names

V1: **CREATE VIEW** WORKS_ON1
 AS SELECT Fname, Lname, Pname, Hours
 FROM EMPLOYEE, PROJECT, WORKS_ON
 WHERE Ssn=Essn **AND** Pno=Pnumber;

V2: **CREATE VIEW** DEPT_INFO(Dept_name, No_of_emps, Total_sal)
 AS SELECT Dname, **COUNT** (*), **SUM** (Salary)
 FROM DEPARTMENT, EMPLOYEE
 WHERE Dnumber=Dno
 GROUP BY Dname;

SUBQUERY

- A *subquery* is a query within a query.
- Subqueries enable you to write queries that select data rows for criteria that are actually developed while the query is executing at *run time*.

Example

```
SELECT emp_last_name "Last Name", emp_first_name "First  
Name",  
       emp_salary "Salary"  
FROM employee  
WHERE emp_salary =  
      (SELECT MIN(emp_salary)  
       FROM employee);
```

Last Name	First Name	Salary
Markis	Marcia	\$25,000
Amin	Hyder	\$25,000
Prescott	Sherri	\$25,000

SUBQUERY TYPES

- by use of the IN operator or with a comparison operator modified by the ANY or ALL optional keywords.
- These subqueries can return a group of values, but the values must be from a single column of a table.

SUBQUERY TYPES

2. Subqueries that use an unmodified comparison operator (=, <, >, <>) – these subqueries must return only a single, *scalar* value.
3. Subqueries that use the EXISTS operator to test the *existence* of data rows satisfying specified criteria.

SUBQUERY – General Rules

A subquery SELECT statement is very similar to the SELECT statement used to begin a regular or outer query. The complete syntax of a subquery is shown below.

```
( SELECT subquery_select_argument
  FROM {table_name | view_name}
       {table_name | view_name} ...
  [WHERE search_conditions]
  [GROUP BY aggregate_expression [, aggregate_expression] ...]
  [HAVING search_conditions] )
```

Rules Cont'd

- The SELECT clause of a subquery must contain only one expression, only one aggregate function, or only one column name.
- The value(s) returned by a subquery must be *join-compatible* with the WHERE clause of the outer query.

Example

```
SELECT emp_last_name "Last Name",  
       emp_first_name "First Name"  
FROM employee  
WHERE emp_ssn IN  
       (SELECT dep_emp_ssn  
        FROM dependent);
```

Last Name	First Name
-----------	------------

Bock	Douglas
Zhu	Waiman
Joyner	Suzanne

Rules Cont'd

Subqueries cannot manipulate their results internally. This means that a subquery cannot include the ORDER BY clause, the COMPUTE clause, or the INTO keyword.

SUBQUERIES AND THE IN Operator

- Subqueries that are introduced with the keyword **IN** take the general form:
 - WHERE expression [NOT] IN (subquery)

```
SELECT emp_last_name "Last Name",  
       emp_first_name "First Name"  
FROM employee  
WHERE emp_ssn IN  
      (SELECT dep_emp_ssn  
       FROM dependent  
       WHERE dep_gender = 'M');
```

Last Name	First Name
Bock	Douglas
Zhu	Waiman
Joyner	Suzanne

SUBQUERIES AND THE IN Operator

- Conceptually, this statement is evaluated in two steps.
- First, the inner query returns the identification numbers of those employees that have male dependents.

```
SELECT dep_emp_ssn
FROM dependent
WHERE dep_gender = 'M';
DEP_EMP_S
-----
999444444
999555555
999111111
```

SUBQUERIES AND COMPARISON OPERATORS

- The general form of the WHERE clause with a comparison operator is similar to that used thus far in the text.
- Note that the subquery is again enclosed by parentheses.

WHERE <expression> <comparison_operator> (subquery)

SUBQUERIES AND COMPARISON OPERATORS

- The most important point to remember when using a subquery with a comparison operator is that the subquery can only return a single or *scalar* value.
- This is also termed a *scalar subquery* because a single column of a single row is returned by the subquery.
- If a subquery returns more than one value, the Oracle Server will generate the “ORA-01427: *single-row subquery returns more than one row*” error message, and the query will fail to execute.

SUBQUERIES AND COMPARISON OPERATORS

- Let's examine a subquery that will not execute because it violates the "single value" rule.
- The query shown below returns multiple values for the *emp_salary* column.

```
SELECT emp_salary  
FROM employee  
WHERE emp_salary > 40000;
```

```
EMP_SALARY
```

```
-----
```

```
55000
```

```
43000
```

```
43000
```

SUBQUERIES AND COMPARISON OPERATORS

- If we substitute this query as a subquery in another SELECT statement, then that SELECT statement will fail.
- This is demonstrated in the next SELECT statement. Here the SQL code will fail because the subquery uses the greater than (>) comparison operator and the subquery returns multiple values.

```
SELECT emp_ssn  
FROM employee  
WHERE emp_salary >  
      (SELECT emp_salary  
       FROM employee  
        WHERE emp_salary > 40000);
```

ERROR at line 4:

ORA-01427: single-row subquery returns more than one row

Aggregate Functions and Comparison Operators

- The aggregate functions (AVG, SUM, MAX, MIN, and COUNT) always return a *scalar* result table.
- Thus, a subquery with an aggregate function as the object of a comparison operator will always execute provided you have formulated the query properly.

Aggregate Functions and Comparison Operators

```
SELECT emp_last_name "Last Name",  
       emp_first_name "First Name",  
       emp_salary "Salary"  
FROM employee  
WHERE emp_salary >  
      (SELECT AVG(emp_salary)  
       FROM employee);
```

Last Name	First Name	Salary
Bordoloi	Bijoy	\$55,000
Joyner	Suzanne	\$43,000
Zhu	Waiman	\$43,000
Joshi	Dinesh	\$38,000

Comparison Operators Modified with the ALL or ANY Keywords

- The ALL and ANY keywords can modify a comparison operator to allow an outer query to accept multiple values from a subquery.
- The general form of the WHERE clause for this type of query is shown here.

WHERE <expression> <comparison_operator> [ALL | ANY]
(subquery)

- Subqueries that use these keywords may also include GROUP BY and HAVING clauses.

The ALL Keyword

- The ALL keyword modifies the greater than comparison operator to mean greater than all values.

```
SELECT emp_last_name "Last Name",  
       emp_first_name "First Name",  
       emp_salary "Salary"  
FROM employee  
WHERE emp_salary > ALL  
      (SELECT emp_salary  
       FROM employee  
       WHERE emp_dpt_number = 7);
```

Last Name	First Name	Salary

Bordoloi	Bijoy	\$55,000

The ANY Keyword

- The ANY keyword is not as restrictive as the ALL keyword.
- When used with the greater than comparison operator, "> ANY" means greater than some value.

```
SELECT emp_last_name "Last Name",  
       emp_first_name "First Name",  
       emp_salary "Salary"  
FROM employee  
WHERE emp_salary > ANY  
      (SELECT emp_salary  
       FROM employee  
       WHERE emp_salary > 30000);
```

Last Name	First Name	Salary
Bordoloi	Bijoy	\$55,000
Joyner	Suzanne	\$43,000
Zhu	Waiman	\$43,000

Subqueries and the EXISTS operator

- When a subquery uses the EXISTS operator, the subquery functions as an *existence test*.
- The WHERE clause of the outer query tests for the existence of rows returned by the inner query.
- The subquery does not actually produce any data; rather, it returns a value of TRUE or FALSE.

Subqueries and the EXISTS operator

- The general format of a subquery WHERE clause with an EXISTS operator is shown here.
- Note that the NOT operator can also be used to negate the result of the EXISTS operator.

WHERE [NOT] EXISTS (subquery)

Example

```
SELECT emp_last_name "Last Name", emp_first_name "First  
Name"  
FROM employee  
WHERE EXISTS  
    (SELECT *  
     FROM dependent  
     WHERE emp_ssn = dep_emp_ssn);
```

Last Name	First Name
Joyner	Suzanne
Zhu	Waiman
Bock	Douglas

Subqueries and the EXISTS operator

- The EXISTS operator is very important, because there is often no alternative to its use.
- All queries that use the IN operator or a modified comparison operator (=, <, >, etc. modified by ANY or ALL) can be expressed with the EXISTS operator.
- However, some queries formulated with EXISTS cannot be expressed in any other way!

Subqueries and the EXISTS operator

```
SELECT emp_last_name
FROM employee
WHERE emp_ssn = ANY
      (SELECT dep_emp_ssn
       FROM dependent);
```

EMP_LAST_NAME

Bock

Zhu

Joyner

```
SELECT
      emp_last_name
FROM employee
WHERE EXISTS
      (SELECT *
       FROM dependent
        WHERE emp_ssn
              = dep_emp_ssn);
```

EMP_LAST_NAME

Bock

Zhu

Joyner

Subqueries and the EXISTS operator

The NOT EXISTS operator is the mirror-image of the EXISTS operator. A query that uses NOT EXISTS in the WHERE clause is satisfied if the subquery returns no rows.

VIEWS

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
100	Steven	King	SKING	515.123.4567	17-JUN-87	AD_FRES	2400
101	Neena	Kochhar	NKOCHHAR	515.123.4568	21-SEP-89	AD_VP	1700
102	Lex	De Haan	LDEHAAN	515.123.4569	13-JAN-93	AD_VP	1700
103	Alexander	Hunold	AHUNOLD	590.423.4567	03-JAN-90	IT_PROG	900
104	Bruce	Ernst	BERNST	590.423.4568	21-MAY-91	IT_PROG	600
107	Diana	Lorentz	DLORENTZ	590.423.5567	07-FEB-98	IT_PROG	420
124	Ivy	Mourgos	IMOURGOS	650.123.5234	16-NOV-89	ST_MAN	580
141	Trenna	Ras	TRAS	650.121.3009	17-OCT-95	ST_CLERK	350
142	Curis	Davies	CDAVIES	650.121.2994	29-JAN-97	ST_CLERK	310
143	Randall	Mateo	RMATEO	650.121.2074	15-MAR-90	ST_CLERK	260
149	Zlotkey				29-JUL-96	ST_CLERK	250
174	Abel				24-JAN-00	SA_MAN	1050
176	Taylor				11-MAY-96	SA_REP	1100
176	Taylor				11-MAR-98	SA_REP	860
176	Kimberly	Grant	KGRANT	611.44.1044.429203	24-MAY-99	SA_REP	700
200	Jennifer	Whalen	JWHALEN	515.123.4444	17-SEP-87	AD_ASST	440
201	Michael	Hartstein	MHARTSTE	515.123.5555	17-FEB-96	MK_MAN	1300
202	Pat	Fay	PFAY	603.123.6666	17-AUG-97	MK_REP	600
205	Shelley	Higgins	SHIGGINS	515.123.8080	07-JUN-94	AC_MGR	1200
206	William	Gietz	WGIEZT	515.123.8181	07-JUN-94	AC_ACCOUNT	830

20 rows selected.

Why Use Views?

- To restrict data access
- To make complex queries easy
- To provide data independence
- To present different views of the same data

Creating a View

- You embed a subquery within the CREATE VIEW statement.

```
CREATE [OR REPLACE] [FORCE|NOFORCE] VIEW view  
    [(alias[, alias]...)]  
    AS subquery  
[WITH CHECK OPTION [CONSTRAINT constraint]]  
[WITH READ ONLY [CONSTRAINT constraint]];
```

Creating a View

- Create a view, EMPVU80, that contains details of employees in department 80.

```
CREATE VIEW empvu80
AS SELECT employee_id, last_name, salary
FROM employees
WHERE department_id = 80;
```

View created.

- Describe the structure of the view by using the `SQL*Plus` `DESCRIBE` command.

```
DESCRIBE empvu80
```

Creating a View

- Create a view by using column aliases in the subquery.

```
CREATE VIEW  salvu50
  AS SELECT  employee_id ID_NUMBER, last_name NAME,
            salary*12 ANN_SALARY
    FROM      employees
   WHERE      department_id = 50;
```

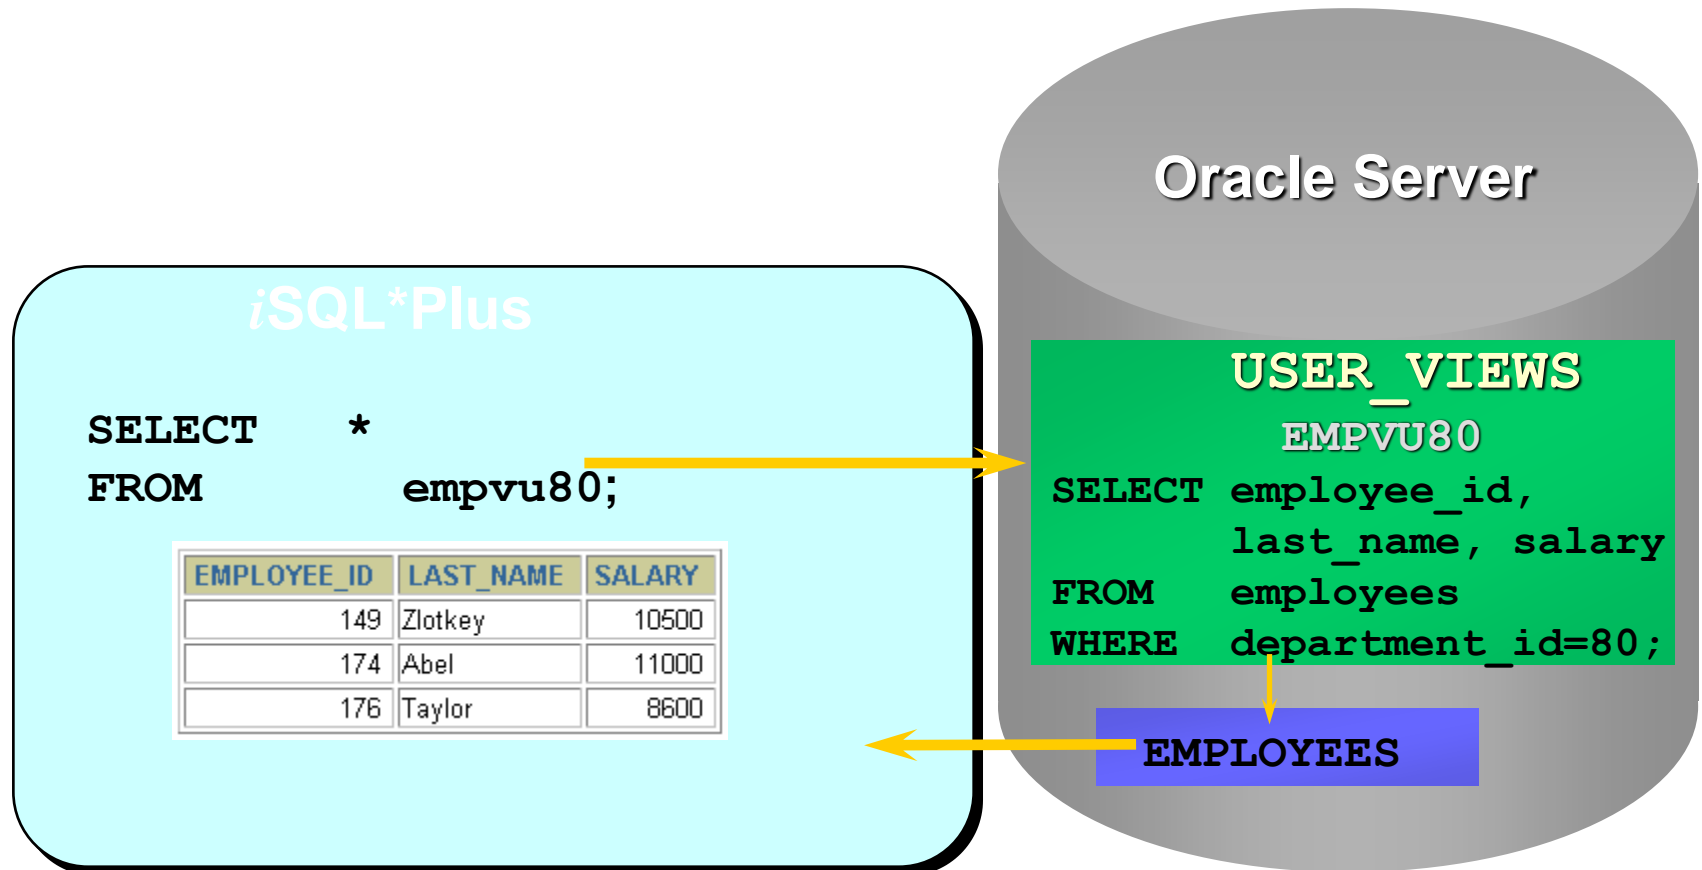
- **View created.** Select the columns from this view by the given alias names.

Retrieving Data from a View

```
SELECT *  
FROM salvu50;
```

ID_NUMBER	NAME	ANN_SALARY
124	Mourgos	69600
141	Rajs	42000
142	Davies	37200
143	Matos	31200
144	Vargas	30000

Querying a View



Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

```
CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]
                             { , <column name> <column type> [ <attribute constraint> ] }
                             [ <table constraint> { , <table constraint> } ] )
```

```
DROP TABLE <table name>
ALTER TABLE <table name> ADD <column name> <column type>
```

```
SELECT [ DISTINCT ] <attribute list>
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }
[ WHERE <condition> ]
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) )
                      { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) } ) )
```

```
<grouping attributes> ::= <column name> { , <column name> }
```

```
<order> ::= ( ASC | DESC )
```

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }
| <select statement> )
```

continued on next slide

Summary of SQL Syntax

Table 7.2 Summary of SQL Syntax

DELETE FROM <table name>

[WHERE <selection condition>]

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[WHERE <selection condition>]

CREATE [UNIQUE] INDEX <index name>

ON <table name> (<column name> [<order>] { , <column name> [<order>] })

[CLUSTER]

DROP INDEX <index name>

CREATE VIEW <view name> [(<column name> { , <column name> })]

AS <select statement>

DROP VIEW <view name>

NOTE: The commands for creating and dropping indexes are not part of standard SQL.