



# **Database Recovery Techniques**

**Dr. Sumeet Kaur Sehra**

# Introduction

- Recovery algorithms
- Recovery concepts
  - Write-ahead logging
  - In-place versus shadow updates
  - Rollback
  - Deferred update
  - Immediate update
- Certain recovery techniques best used with specific concurrency control methods

# Recovery Concepts

- Recovery process restores database to most recent consistent state before time of failure
- Information kept in system log
- Typical recovery strategies
  - Restore backed-up copy of database
    - Best in cases of extensive damage
  - Identify any changes that may cause inconsistency
    - Best in cases of noncatastrophic failure
    - Some operations may require redo

# Recovery Concepts (cont'd.)

## Deferred update techniques

- Do not physically update the database until after transaction commits
- Undo is not needed; redo may be needed

## Immediate update techniques

- Database may be updated by some operations of a transaction before it reaches commit point
- Operations also recorded in log
- Recovery still possible

# Recovery Concepts (cont'd.)

## Undo and redo operations required to be idempotent

- Executing operations multiple times equivalent to executing just once
- Entire recovery process should be idempotent

## Caching (buffering) of disk blocks

- DBMS cache: a collection of in-memory buffers
- Cache directory keeps track of which database items are in the buffers

## Recovery Concepts (cont'd.)

Cache buffers replaced (flushed) to make space for new items

Dirty bit associated with each buffer in the cache

- Indicates whether the buffer has been modified

Contents written back to disk before flush if dirty bit equals one

Pin-unpin bit

- Page is pinned if it cannot be written back to disk yet

# Recovery Concepts (cont'd.)

## Main strategies

- In-place updating
  - Writes the buffer to the same original disk location
  - Overwrites old values of any changed data items
- Shadowing
  - Writes an updated buffer at a different disk location, to maintain multiple versions of data items
  - Not typically used in practice

Before-image: old value of data item

After-image: new value of data item

# Recovery Concepts (cont'd.)



## Write-ahead logging

Ensure the before-image (BFIM) is recorded

Appropriate log entry flushed to disk

Necessary for UNDO operation if needed



## UNDO-type log entries



## REDO-type log entries



# Recovery Concepts (cont'd.)



Steal/no-steal and  
force/no-force

Specify rules that govern when a page from the database cache can be written to disk



No-steal approach

Cache buffer page updated by a transaction cannot be written to disk before the transaction commits



Steal approach

Recovery protocol allows writing an updated buffer before the transaction commits

# Recovery Concepts (cont'd.)



## **Force approach**

All pages updated by a transaction  
are immediately written to disk  
before the transaction commits


Otherwise, no-force



## **Typical database systems employ a steal/no-force strategy**

Avoids need for very large buffer  
space

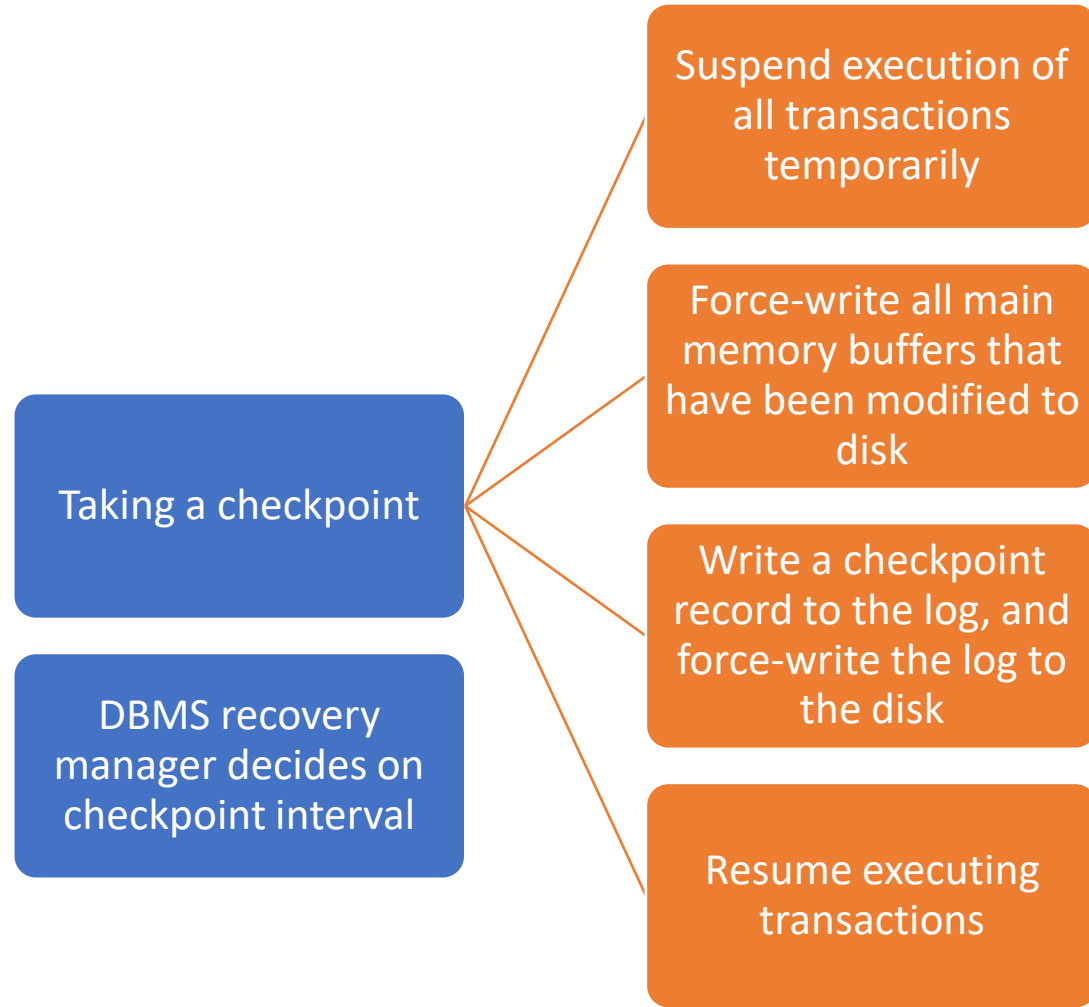
Reduces disk I/O operations for  
heavily updated pages



# Recovery Concepts (cont'd.)

- Write-ahead logging protocol for recovery algorithm requiring both UNDO and REDO
  - BFIM of an item cannot be overwritten by its after image until all UNDO-type log entries have been force-written to disk
  - Commit operation of a transaction cannot be completed until all REDO-type and UNDO-type log records for that transaction have been force-written to disk

# Checkpoints in the System Log and Fuzzy Checkpointing



# Checkpoints in the System Log and Fuzzy Checkpointing (cont'd.)

- Fuzzy checkpointing
  - System can resume transaction processing after a begin\_checkpoint record is written to the log
  - Previous checkpoint record maintained until end\_checkpoint record is written

# Transaction Rollback

## Transaction failure after update but before commit

- Necessary to roll back the transaction
- Old data values restored using undo-type log entries

## Cascading rollback

- If transaction T is rolled back, any transaction S that has read value of item written by T must also be rolled back
- Almost all recovery mechanisms designed to avoid this

Illustrating cascading  
rollback (a process that  
never occurs  
in strict or cascadeless  
schedules) (a) The read and  
write operations of three  
transactions (b) System log  
at point of crash (c)  
Operations before the crash

(a)

$T_1$	$T_2$	$T_3$
read_item(A)	read_item(B)	read_item(C)
read_item(D)	write_item(B)	write_item(B)
write_item(D)	read_item(D)	read_item(A)
	write_item(D)	write_item(A)

(b)

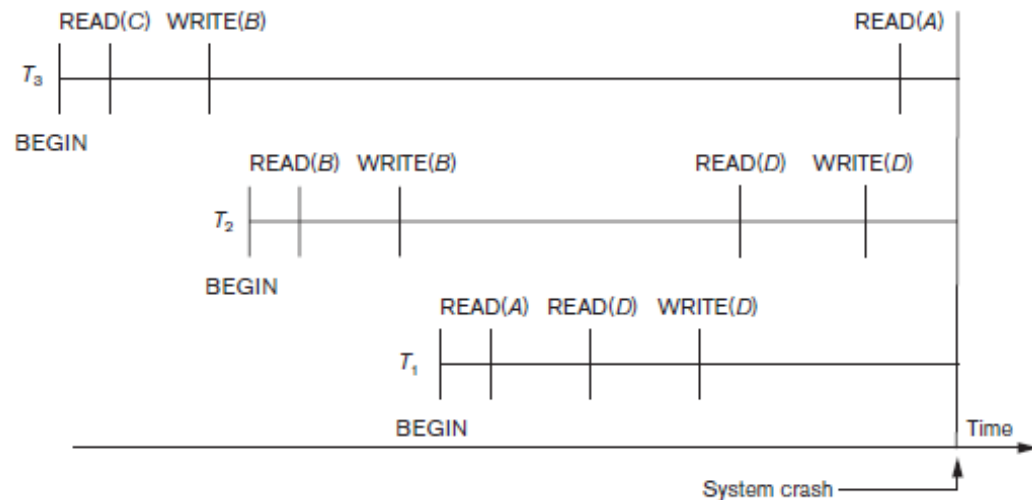
	A	B	C	D
	30	15	40	20
[start_transaction, $T_3$ ]				
[read_item, $T_3$ , C]				
* [write_item, $T_3$ , B, 15, 12]		12		
[start_transaction, $T_2$ ]				
[read_item, $T_2$ , B]				
** [write_item, $T_2$ , B, 12, 18]		18		
[start_transaction, $T_1$ ]				
[read_item, $T_1$ , A]				
[read_item, $T_1$ , D]				
[write_item, $T_1$ , D, 20, 25]				25
[read_item, $T_2$ , D]				
** [write_item, $T_2$ , D, 25, 26]				26
[read_item, $T_3$ , A]				

← System crash

\*  $T_3$  is rolled back because it  
did not reach its commit point.

\*\*  $T_2$  is rolled back because it  
reads the value of item B written by  $T_3$ .

(c)



# Transactions that Do Not Affect the Database

Example actions: generating and printing messages and reports

If transaction fails before completion, may not want user to get these reports

- Reports should be generated only after transaction reaches commit point

Commands that generate reports issued as batch jobs executed only after transaction reaches commit point

- Batch jobs canceled if transaction fails



# NO-UNDO/REDO Recovery Based on Deferred Update

---

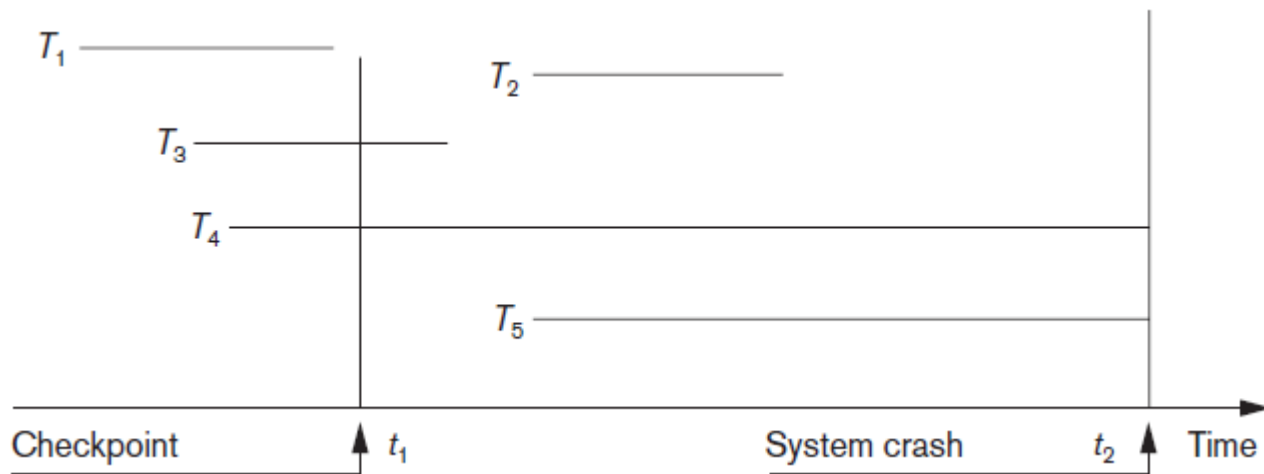
- Deferred update concept
  - Postpone updates to the database on disk until the transaction completes successfully and reaches its commit point
  - Redo-type log entries are needed
  - Undo-type log entries not necessary
  - Can only be used for short transactions and transactions that change few items
    - Buffer space an issue with longer transactions

# NO-UNDO/REDO Recovery Based on Deferred Update (cont'd.)

---

- Deferred update protocol
  - Transaction cannot change the database on disk until it reaches its commit point
    - All buffers changed by the transaction must be pinned until the transaction commits (no-steal policy)
  - Transaction does not reach its commit point until all its REDO-type log entries are recorded in log and log buffer is force-written to disk

# NO-UNDO/REDO Recovery Based on Deferred Update (cont'd.)



An example of a recovery timeline to illustrate the effect of checkpointing

# Recovery Techniques Based on Immediate Update



**Database can be updated immediately**

No need to wait for transaction to reach commit point

Not a requirement that every update be immediate



**UNDO-type log entries must be stored**



**Recovery algorithms**

UNDO/NO-REDO (steal/force strategy)

UNDO/REDO (steal/no-force strategy)

$T_1$
read_item(A)
read_item(D)
write_item(D)

$T_2$
read_item(B)
write_item(B)
read_item(D)
write_item(D)

$T_3$
read_item(A)
write_item(A)
read_item(C)
write_item(C)

$T_4$
read_item(B)
write_item(B)
read_item(A)
write_item(A)

An example of recovery using deferred update with concurrent transactions (a) The READ and WRITE operations of four transactions (b) System log at the point of crash

[start_transaction, $T_1$ ]
[write_item, $T_1$ , D, 20]
[commit, $T_1$ ]
[checkpoint]
[start_transaction, $T_4$ ]
[write_item, $T_4$ , B, 15]
[write_item, $T_4$ , A, 20]
[commit, $T_4$ ]
[start_transaction, $T_2$ ]
[write_item, $T_2$ , B, 12]
[start_transaction, $T_3$ ]
[write_item, $T_3$ , A, 30]
[write_item, $T_2$ , D, 25]

← System crash

$T_2$  and  $T_3$  are ignored because they did not reach their commit points.

$T_4$  is redone because its commit point is after the last system checkpoint.

## 22.4 Shadow Paging

No log required in a single-user environment

- Log may be needed in a multiuser environment for the concurrency control method

Shadow paging considers disk to be made of  $n$  fixed-size disk pages

- Directory with  $n$  entries is constructed
- When transaction begins executing, directory copied into shadow directory to save while current directory is being used
- Shadow directory is never modified

# Shadow Paging (cont'd.)

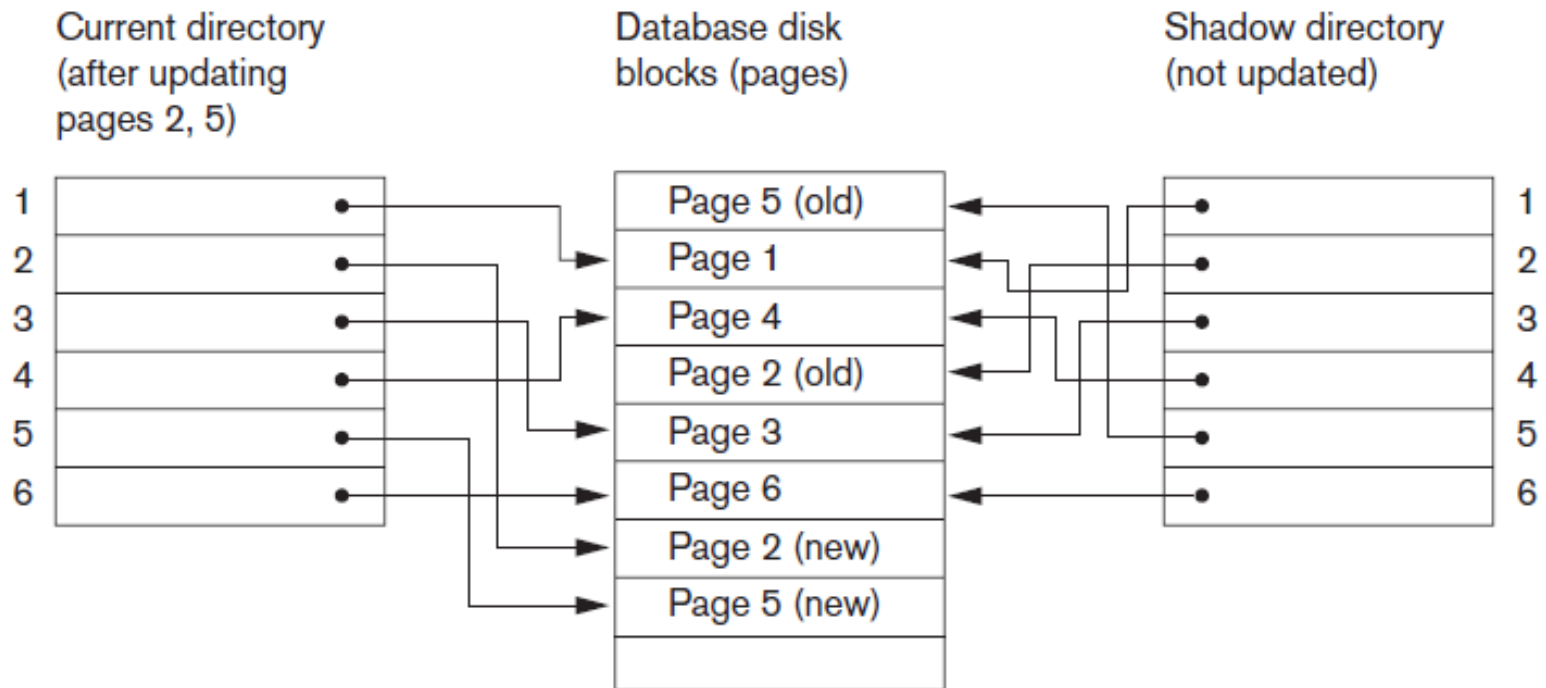
New copy of the modified page created and stored elsewhere

- Current directory modified to point to new disk block
- Shadow directory still points to old disk block

## Failure recovery

- Discard current directory
- Free modified database pages
- NO-UNDO/NO-REDO technique

# Shadow Paging (cont'd.)



<sup>5</sup>The directory is similar to the page table maintained by the operating system for each process.

An example of shadow paging



# Recovery in Multidatabase Systems

Two-level recovery mechanism

Global recovery manager (coordinator) needed to maintain recovery information

Coordinator follows two-phase commit protocol

- Phase 1: Prepare for commit message
  - Ready to commit or cannot commit signal returned
- Phase 2: Issue commit signal

Either all participating databases commit the effect of the transaction or none of them do

# Recovery in Multidatabase Systems (cont'd.)

Always possible to recover to a state where either the transaction is committed or it is rolled back

Failure during phase 1 requires rollback

Failure during phase 2 means successful transaction can recover and commit

# Database Backup and Recovery from Catastrophic Failures

## Database backup

- Entire database and log periodically copied onto inexpensive storage medium
- Latest backup copy can be reloaded from disk in case of catastrophic failure

## Backups often moved to physically separate locations

- Subterranean storage vaults

## Database Backup and Recovery from Catastrophic Failures (cont'd.)

- Backup system log at more frequent intervals and copy to magnetic tape
  - System log smaller than database
    - Can be backed up more frequently
- Benefit: users do not lose all transactions since last database backup