

Operating Systems

The Abstraction: The Process

These slides include content from the work of:
Youjip Won, KIAT OS Lab

Operating System

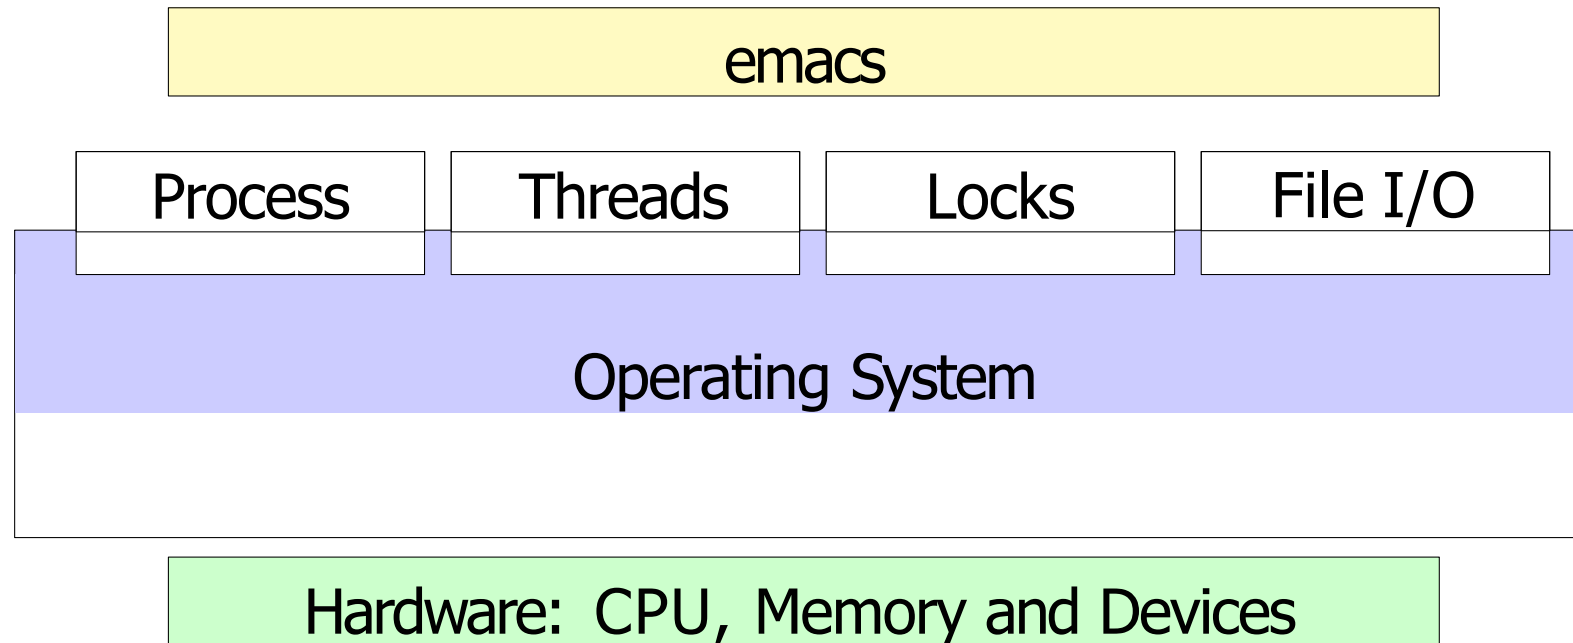
```
graph TD; emacs[emacs] --- OS[Operating System]; OS --- hardware[Hardware: CPU, Memory and Devices];
```

emacs

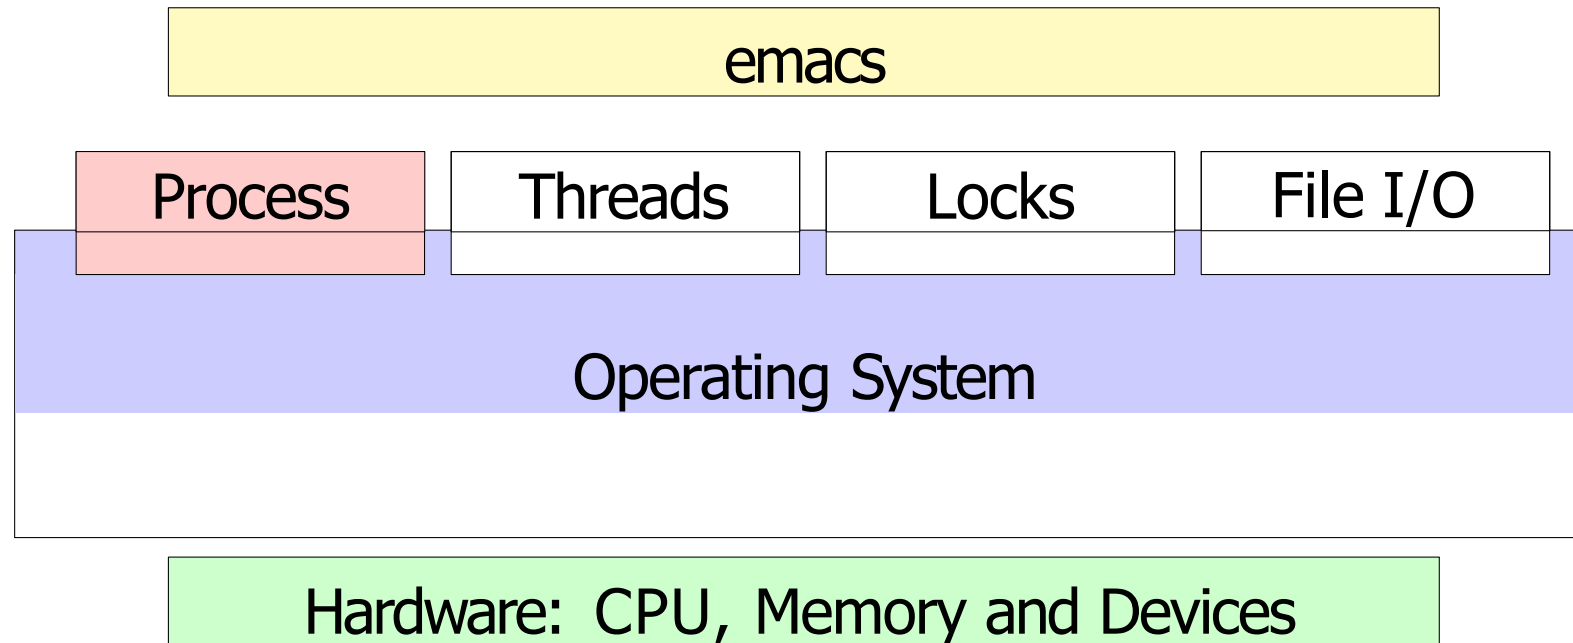
Operating System

Hardware: CPU, Memory and Devices

Operating System: Basic Abstractions and APIs



Today: Introduce the Process Abstraction



How to provide the illusion of many CPUs?

▣ CPU virtualizing


- ◆ The OS can promote the illusion that many virtual CPUs exist.
- ◆ **Time sharing**: Running one process, then stopping it and running another
 - The potential cost is **performance**: as each will run more slowly if the CPU(s) must be shared.

Processes


- A process is an instance of a program running
- Examples (can all run simultaneously):
 - ▶ `gcc file_A.c` – compiler running on file A
 - ▶ `gcc file_B.c` – compiler running on file B
 - ▶ `emacs` – text editor
 - ▶ `firefox` – web browser
- Non-examples (implemented as one process):
 - Multiple older versions of Firefox (still one process)
- Modern OSes run multiple processes simultaneously

Speed


- Multiple processes can increase CPU utilization
 - ▶ Overlap one process's computation with another's wait



- Multiple processes can reduce latency
 - ▶ Running *A* then *B* requires 100 sec for *B* to complete



- ▶ Running *A* and *B* concurrently makes *B* finish faster




A process is a **running program**.

□ Process comprising of:

1. Memory:

- ♦ **Text Segment (Instructions):** This is the memory space where the executable code of the process is loaded.
- ♦ **Data Segment:** This holds the global and static variables used by the process.
- ♦ **Heap:** This is the dynamically allocated portion of memory that grows and shrinks as the process makes system calls like `malloc` and `free` in C, for instance.
- ♦ **Stack:** This memory is used for function call management, local variables, return addresses

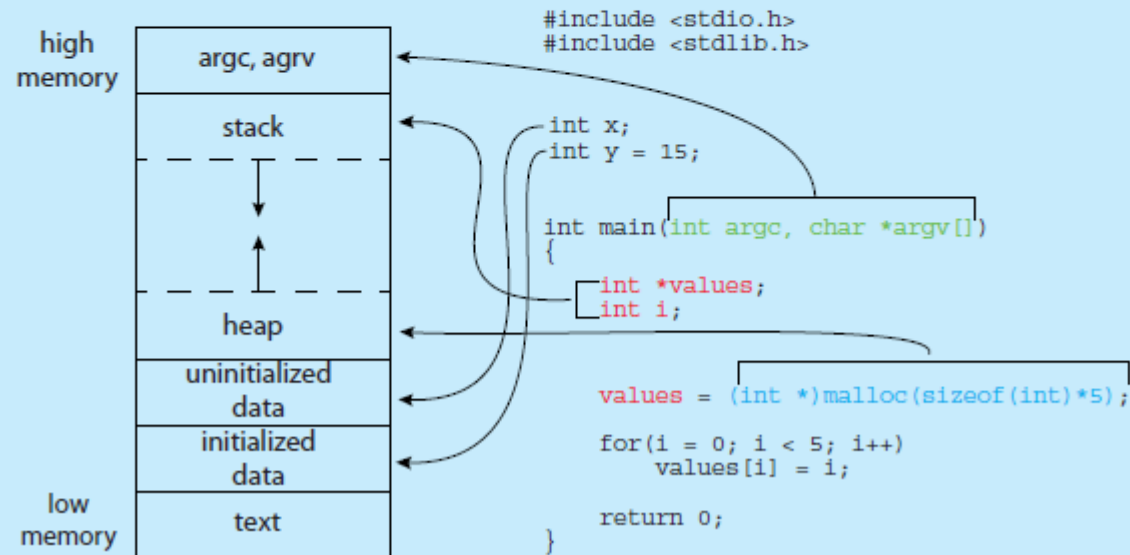
2. Process Control Block (PCB):



Memory Layout of a C Program

The figure shown below illustrates the layout of a C program in memory, highlighting how the different sections of a process relate to an actual C program. This figure is similar to the general concept of a process in memory as shown in Figure 3.1, with a few differences:

- The global data section is divided into different sections for (a) initialized data and (b) uninitialized data.
- A separate section is provided for the `argc` and `argv` parameters passed to the `main()` function.



Process Control Block (PCB)

- OS keeps data structure for each proc
 - ▶ Process Control Block (PCB)
 - ▶ Called `proc` in Unix, `task_struct` in Linux, and `struct Process` in COS
- Process ID (PID): A unique identifier for the process.
- Process State: The current state (e.g., ready, running, waiting, terminated).
- Program Counter: The address of the next instruction to execute.
- CPU Registers: The values of all CPU registers for this process.
- I/O Status Information: List of I/O devices allocated to the process, list of open files, etc.
-

Process state
Process ID
User id, etc.
Program counter
Registers
Address space (VM data structs)
Open files

PCB

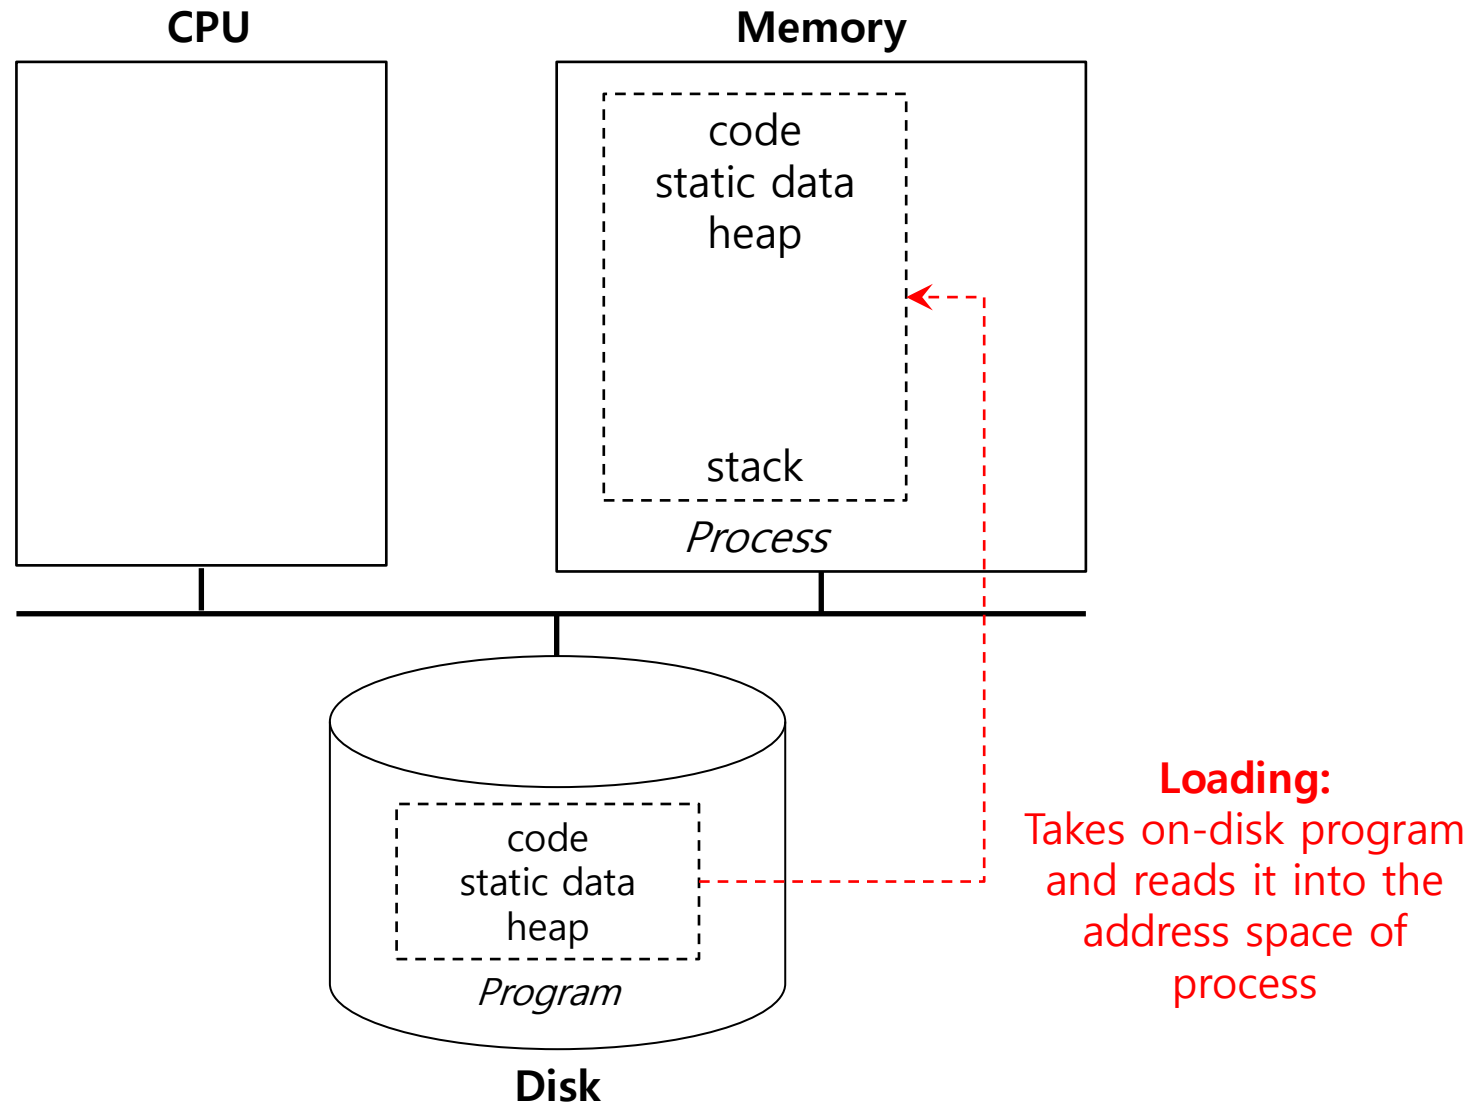
Process Creation

1. **Load** a program code into memory, into the address space of the process.
 - ◆ Programs initially reside on disk in *executable format*.
 - ◆ In early (or simple) operating systems, the loading process is done **eagerly**.
 - all at once before running the program.
 - ◆ modern OSes perform the loading process **lazily**.
 - Loading pieces of code or data only as they are needed during program execution.
2. The program's run-time **stack** is allocated.
 - ◆ Use the stack for *local variables*, *function parameters*, and *return address*.
 - ◆ Initialize the stack with arguments → `argc` and the `argv` array of `main()` function

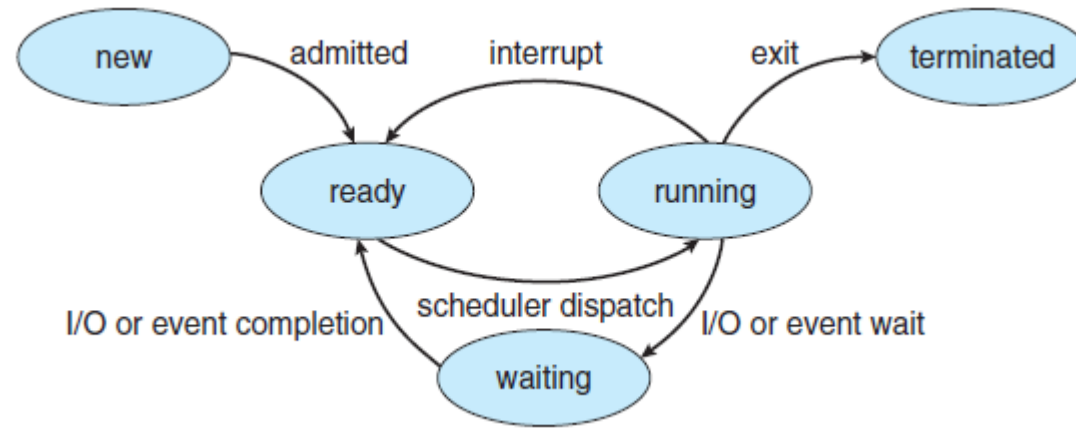
Process Creation (Cont.)

3. The program's **heap** is created.
 - ◆ Used for explicitly requested dynamically allocated data.
 - ◆ Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS do some other initialization tasks.
 - ◆ input/output (I/O) setup
5. **Start the program** running at the entry point, namely `main()`.
 - ◆ The OS *transfers control* of the CPU to the newly-created process.

Loading: From Program To Process



Process State Transition



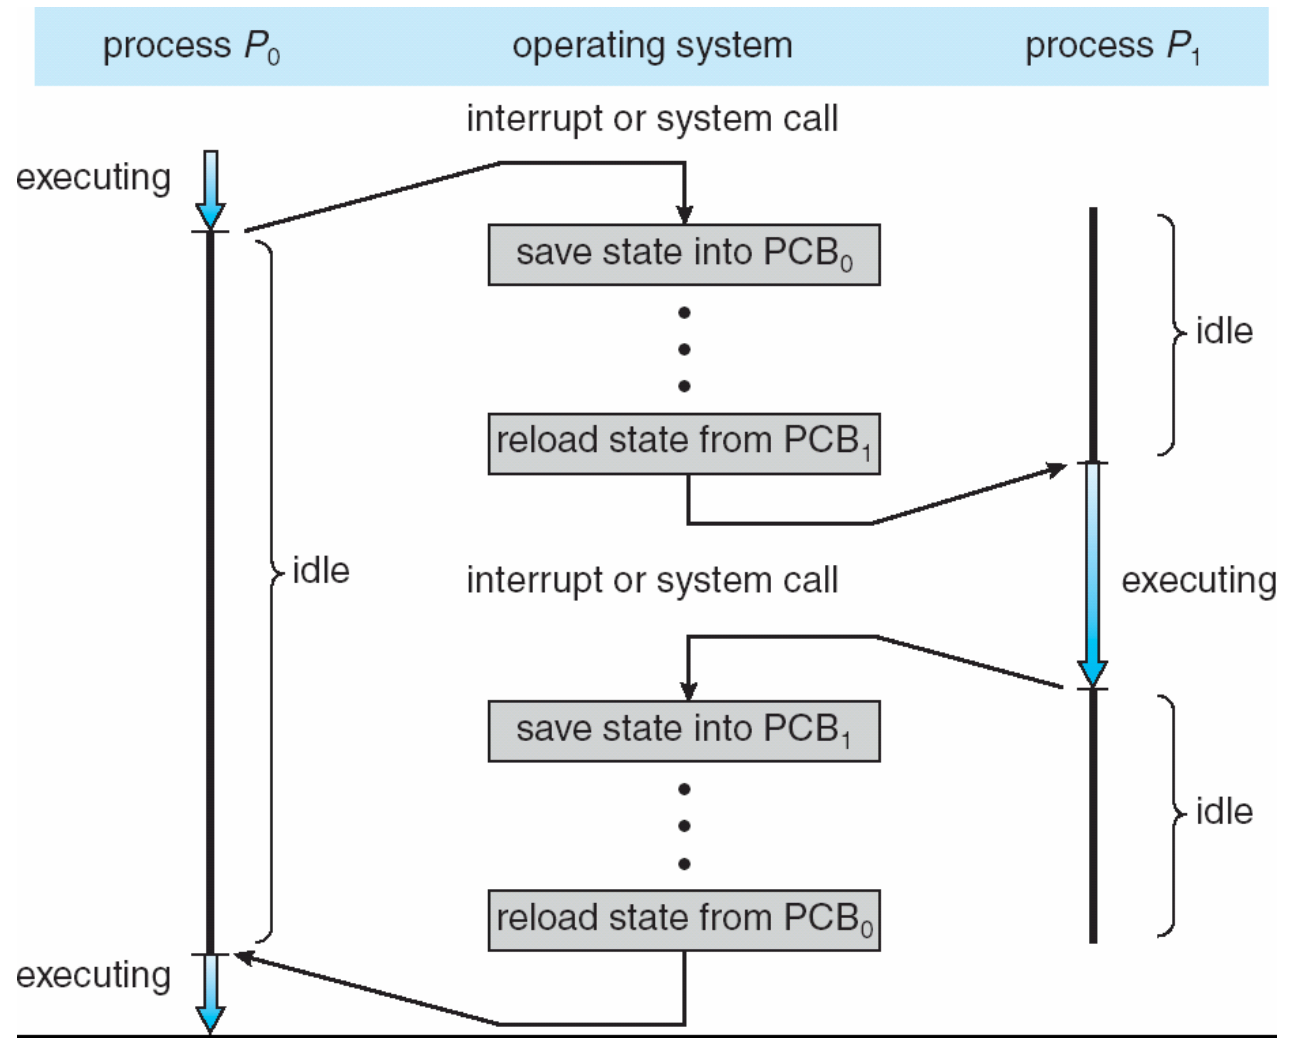
- ▣ Process can be in one of several states
 - ◆ new & terminated at beginning & end of life
 - ◆ running – currently executing (or will execute on kernel return)
 - ◆ ready – can run, but kernel has chosen different process to run
 - ◆ waiting – needs async event (e.g., disk operation) to proceed

Scheduling

- ▣ An I/O-bound process is one that spends more of its time doing I/O than it spends doing computations.
- ▣ A CPU-bound process, in contrast, generates I/O requests infrequently, using more of its time doing computations.
- ▣ OS maintains scheduling queues of processes
 - ◆ Job queue – set of all processes in the system
 - ◆ Ready queue – set of processes ready and waiting to execute
 - ◆ Device queues – set of processes waiting for an I/O device
- ▣ Processes migrate among the various queues

Context switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**
- An **interrupt** is a signal to the CPU that tells it to stop its current execution and immediately handle an event. Examples: A key is pressed, system calls, CPU clock tick forces scheduler to switch processes



Process Life Cycle

1. Program Request

- ◆ User clicks program / enters command.
- ◆ OS gets request to start execution.

2. Process Creation (New → Ready)

- ◆ OS creates a PCB (Process Control Block) with PID & metadata.
- ◆ Virtual memory allocated:
 - Code segment (program instructions)
 - Data segment (globals, statics)
 - Heap (dynamic allocation)
 - Stack (function calls, locals, return addresses)
- ◆ Arguments/environment prepared on the stack.
- ◆ Process placed in the Ready queue.

Process Life Cycle

3. Scheduling & Dispatch (Ready → Running)

- ◆ CPU scheduler picks a process from Ready queue.
- ◆ Dispatcher loads registers, program counter, etc. from PCB (**Context Switch** happens).
- ◆ Process starts execution at its entry point (*main*).

4. Execution & System Calls (Running ↔ Waiting)

- ◆ While running, the process may:
 - Execute user instructions in user mode.
 - Make system calls (switch to kernel mode).
 - Perform I/O → if waiting, moved to Blocked state.
- ◆ When blocked process waits; CPU is given to another Ready process (**Context Switch** happens).

5. Process Completion (Running → Terminated)

- ◆ When finished, process calls exit.
- ◆ OS releases memory (heap, stack, code/data) and PCB removed.