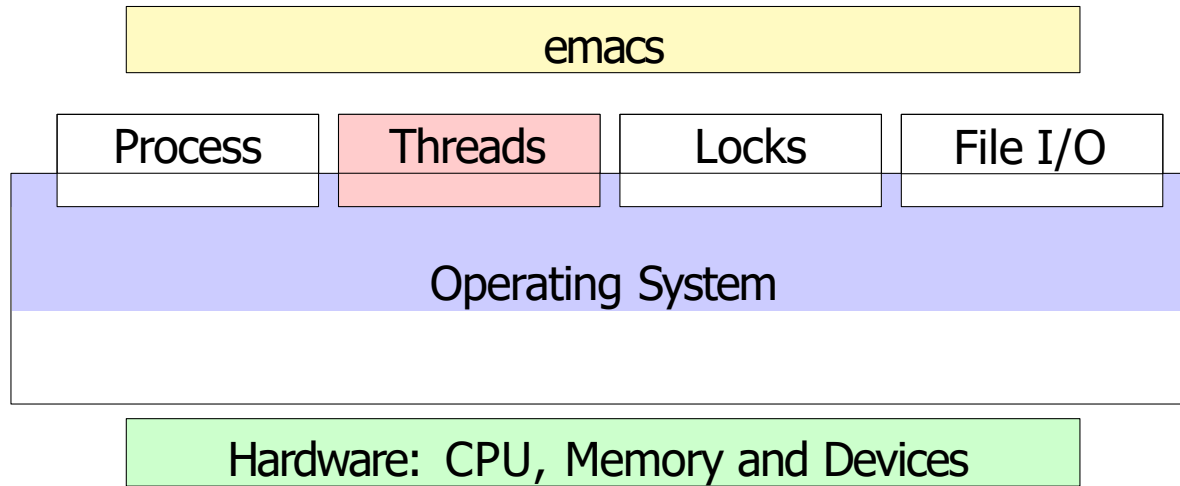


Operating Systems

Threads

These slides include content from the work of:
Youjip Won, KIAT OS Lab

Today: Threads

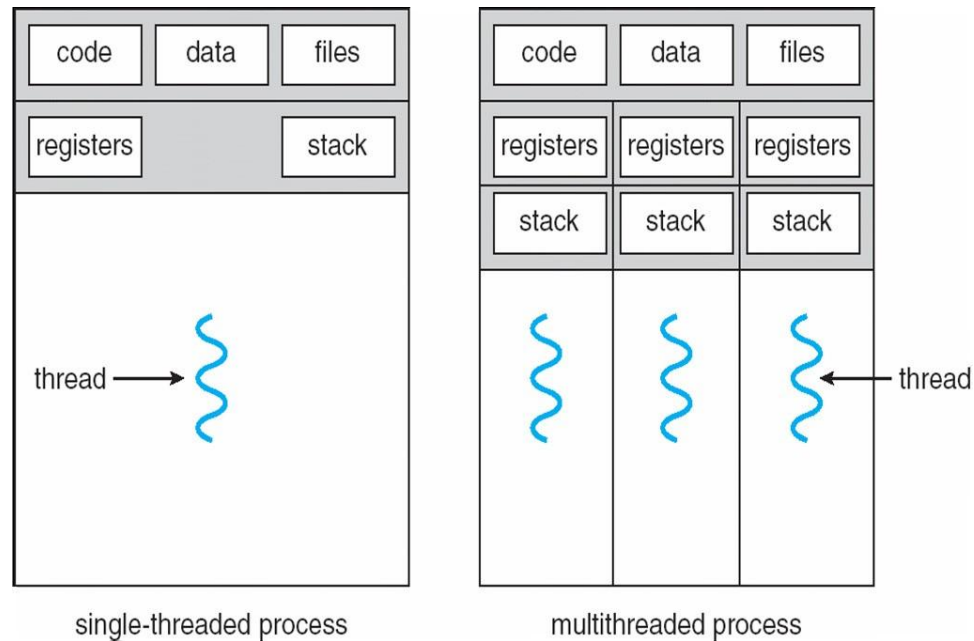


Thread

- ▣ A new abstraction for a single running process
- ▣ A thread is a schedulable execution context
- ▣ It comprises a thread ID, a program counter (PC), a register set, and a stack.

- ▣ Multi-threaded program
 - ◆ A multi-threaded program has more than one point of execution.
 - ◆ Multiple PCs (Program Counter)
 - ◆ They **share** the same **address space**.
 - ◆ If a process has multiple threads of control, it can perform more than one task at a time.

Threads



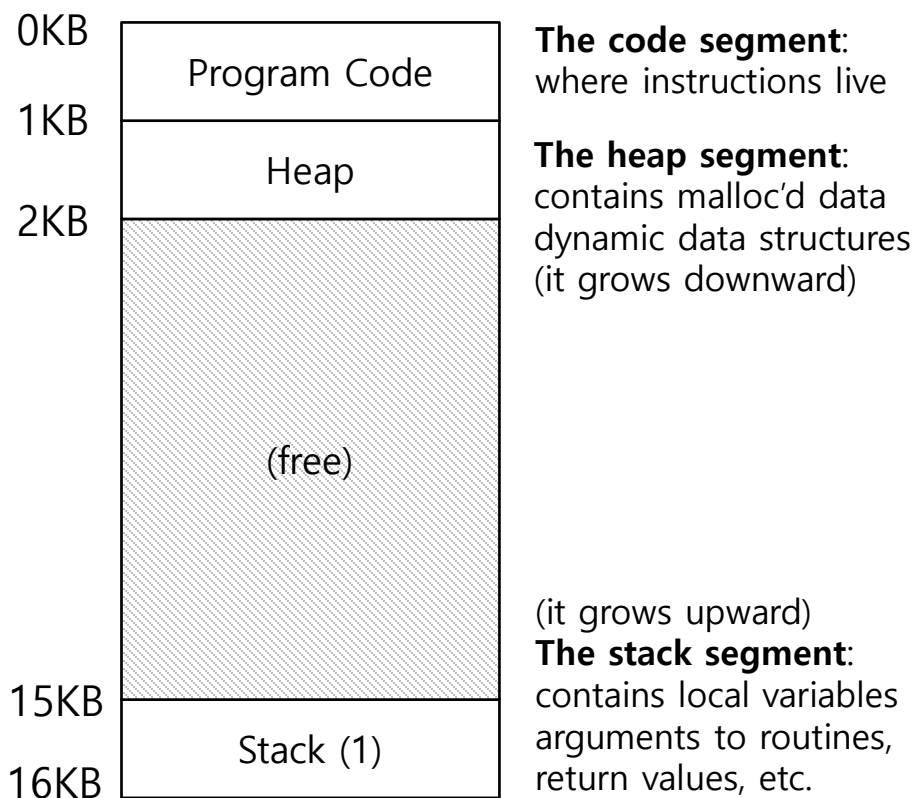
- A thread is a schedulable execution context
 - ▶ Program counter, registers, stack (local variables) . . .
- Multi-threaded programs share the address space (global variables, heap, . . .)

Context switch between threads


- ▣ Each thread has its own program counter and set of registers.
 - ◆ One or more **thread control blocks(TCBs)** are needed to store the state of each thread.
- ▣ When switching from running one (T1) to running the other (T2),
 - ◆ The register state of T1 be saved.
 - ◆ The register state of T2 restored.
 - ◆ The **address space remains** the same.
- ▣ With processes, we saved state to a PCB; now, we'll need TCBs to store the state of each thread of a process.

The stack of the relevant thread

- There will be **one stack per thread**.



**A Single-Threaded
Address Space**



**Two threaded
Address Space**

Why Use Threads?


- ▣ Most popular abstraction for concurrency
 - ◆ Lighter-weight abstraction than processes: **Shared Resources, Context Switching,...**
 - ◆ All threads in one process share memory, file descriptors, etc.
- ▣ Parallelism
 - ◆ Single-threaded program: the task is straightforward, but slow.
 - ◆ Multi-threaded program: natural and typical way to make programs run faster on modern hardware.
 - ◆ **Parallelization**: The task of transforming standard **single-threaded** program into a program that does this sort of work on multiple CPUs.
- ▣ Avoid blocking program progress due to slow I/O.
 - ◆ Threading enables **overlap** of I/O with other activities within a single program.
 - ◆ It is much like **multiprogramming** did for processes across programs.

Why Use Threads?


- Of course, in either of the cases mentioned above, you could use multiple processes instead of threads.
- However, threads share an address space and thus make it easy to share data, and hence are a natural choice when constructing these types of programs.
- Processes are a more sound choice for logically separate tasks where little sharing of data structures in memory is needed.

Concurrency vs Parallelism

Concurrency



Parallelism



Concurrency & Parallelism



Multi-threading Examples

▣ Web Servers:

- ◆ use multi-threading to handle multiple client requests simultaneously.

▣ Web Browsers:

- ◆ use multi-threading to load different elements of a webpage (like text, images, and scripts) at the same time.

▣ Office Applications:

- ◆ Microsoft Word can use multi-threading to perform spell-checking, auto-saving, and complex calculations in the background while still being responsive to the user.

▣ Operating Systems:

- ◆ use multi-threading to manage multiple tasks at once, such as running background services, managing user inputs, and updating the user interface.

▣ Machine Learning Frameworks:

- ◆ Tools like TensorFlow or PyTorch use multi-threading to parallelize data loading and processing.

An Example: Thread Creation

▣ Simple Thread Creation Code (t0.c)

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4  #include "common.h"
5  #include "common_threads.h"
6
7  void *mythread (void *arg) {
8      printf ("%s\n", (char *) arg);
9      return NULL;
10 }
11
12 int main (int argc, char *argv[]) {
13     pthread_t p1, p2;
14     printf("main: begin\n");
15     Pthread_create(&p1, NULL, mythread, "A");
16     Pthread_create(&p2, NULL, mythread, "B");
17     // join waits for their threads to finish
18     Pthread_join(p1, NULL);
19     Pthread_join(p2, NULL);
20     printf("main: end\n");
21     return 0;
22 }
```

Thread Trace (1)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
waits for T1		
<hr/>		
	runs	
	prints "A"	
	returns	
<hr/>		
waits for T2		
<hr/>		
		runs
		prints "B"
		returns
<hr/>		
prints "main: end"		

Thread Trace (2)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
<hr/>		
	runs	
	prints "A"	
	returns	
<hr/>		
creates Thread 2		
<hr/>		
		runs
		prints "B"
		returns
<hr/>		
waits for T1		
returns immediately; T1 is done		
waits for T2		
returns immediately; T2 is done		
prints "main: end"		

Thread Trace (3)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
<hr/>		
		runs
		prints "B"
		returns
<hr/>		
waits for T1		
<hr/>		
	runs	
	prints "A"	
	returns	
<hr/>		
waits for T2		
returns immediately; T2 is done		
prints "main: end"		

Thread Creation

▣ How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*            (*start_routine)(void*),
                    void*            arg);
```

- ◆ `thread`: Used to interact with this thread.
- ◆ `attr`: Used to specify any attributes this thread might have.
 - Stack size, Scheduling priority, ...
- ◆ `start_routine`: the function this thread start running in.
- ◆ `arg`: the argument to be passed to the function (start routine)
 - *a void pointer* allows us to pass in *any type of* argument.

Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- ◆ thread: Specify which thread *to wait for*.
- ◆ value_ptr: A pointer to the return value