

Operating Systems


Scheduling: Introduction

CPU Scheduling


- ❑ The scheduling problem:

- ◆ Have K jobs ready to run
- ◆ Have $N \geq 1$ CPUs
- ◆ Which jobs to assign to which CPU(s)

- ❑ When do we make decision?



CPU Scheduling



- Scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from new/waiting to ready
 4. Switches from ready to running
 5. Exits
- Non-preemptive schedules use 1 & 4 only
- Preemptive schedulers run at all five points

Scheduling Metrics

- ▣ Performance metric: **Turnaround time**
 - ◆ The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- ▣ Throughput – # of process that complete per unit time
- ▣ Response time – time from request to first response
 - ◆ (e.g., key press to character echo, not launch to exit)
- ▣ Above criteria are affected by secondary criteria
 - ◆ CPU utilization – fraction of time CPU doing productive work
 - ◆ Waiting time – time each proc waits in ready queue


Scheduling: Workload assumptions:

- ❑ Workload assumptions:

1. Each job runs for the **same amount of time**.
2. All jobs **arrive** at the same time.
3. All jobs only use the **CPU** (i.e., they perform no I/O).
4. The **run-time** of each job is known.

First In, First Out (FIFO)


- First Come, First Served (FCFS)
 - ◆ Very simple and easy to implement
- Example:
 - ◆ A arrived just before B which arrived just before C (roughly at the same time).
 - ◆ Each job runs for 10 seconds.



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? – Convoy effect


- Let's relax assumption 1: Each job **no longer** runs for the same amount of time.
- Example:
 - A arrived just before B which arrived just before C (roughly at the same time).
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$

Shortest Job First (SJF)


- ❑ Run the shortest job first, then the next shortest, and so on
 - ◆ Non-preemptive scheduler
- ❑ Example:
 - ◆ A arrived just before B which arrived just before C (roughly at the same time).
 - ◆ A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time.
- Example:
 - A arrives at t=0 and needs to run for 100 seconds.
 - B and C arrive at t=10 and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$


Shortest Time-to-Completion First (STCF)

- ▣ Add **preemption** to SJF
 - ◆ Also known as Preemptive Shortest Job First (PSJF)
- ▣ A new job enters the system:
 - ◆ Determine of the remaining jobs and new job
 - ◆ Schedule the job which has the least time left

Shortest Time-to-Completion First (STCF)

Example:

- ◆ A arrives at t=0 and needs to run for 100 seconds.
- ◆ B and C arrive at t=10 and each need to run for 10 seconds




$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$


Example

Process	Arrival Time	Burst Time
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4

- Non-preemptive



- Preemptive



SJF limitations

- Doesn't always minimize average turnaround time
 - ◆ Only minimizes waiting time, which minimizes response time
 - ◆ Example where turnaround time might be suboptimal?
- Can lead to unfairness or starvation
- In practice, can't actually predict the future
- But can estimate CPU burst length based on past

Response time

- The time from **when the job arrives** to the **first time it is scheduled**.

$$T_{response} = T_{firstrun} - T_{arrival}$$

- ◆ STCF and related disciplines are not particularly good for response time.

How can we build a scheduler that is
sensitive to response time?

Round Robin (RR) Scheduling


▣ Time slicing Scheduling

- ◆ Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
 - Time slice is sometimes called a scheduling quantum.
- ◆ It repeatedly does so until the jobs are finished.
- ◆ The length of a time slice must be *a multiple of* the timer-interrupt period.

RR is fair, but performs poorly on metrics such as turnaround time


RR Scheduling Example

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5\ sec$$

SJF (Bad for Response Time)



$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1\ sec$$

RR with a time-slice of 1sec (Good for Response Time)

RR disadvantages

- ❑ Varying sized jobs are good ...what about same-sized jobs?
- ❑ Assume 2 jobs of time=100 each:



- ❑ Even if context switches were free...
 - ◆ What would average completion time be with RR? **199.5**
 - ◆ How does that compare to FCFS? **150**


The length of the time slice is critical.

- ▣ The shorter time slice
 - ◆ Better response time
 - ◆ The cost of context switching will dominate overall performance.

- ▣ The longer time slice
 - ◆ Amortize the cost of switching
 - ◆ Worse response time


Deciding on the length of the time slice presents
a **trade-off** to a system designer

The length of the time slice is critical.



- How to pick quantum?
 - ◆ Want much larger than context switch cost
 - ◆ Majority of bursts should be less than quantum
 - ◆ But not so large system reverts to FCFS
- Typical values: 10–100 msec

Turnaround time vs. quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

Bursts of computation & I/O

- Jobs contain I/O and computation

- Bursts of computation
 - Then must wait for I/O

⋮
⋮

**load store
add store
read from file**

wait for I/O

- To Maximize throughput

- Must maximize CPU utilization
 - Also maximize I/O device utilization

**store increment
index
write to file**

wait for I/O

- How to do?


- Overlap I/O & computation from multiple jobs
 - Means response time very important for I/O-intensive jobs: I/O device will be idle until job gets small amount of CPU to issue next I/O request

**load store
add store
read from file**

wait for I/O

⋮
⋮

Histogram of CPU-burst times




- What does this mean for FCFS?


Incorporating I/O

- ▣ Let's relax assumption 3: The programs perform I/O
- ▣ Example:
 - ◆ A and B need 50ms of CPU time each.
 - ◆ A runs for 10ms and then issues an I/O request
 - I/Os each take 10ms
 - ◆ B simply uses the CPU for 50ms and performs no I/O
 - ◆ The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Poor Use of Resources



Maximize the
CPU utilization

Overlap Allows Better Use of Resources

Incorporating I/O (Cont.)

- ▣ When a job initiates an I/O request.
 - ◆ The job is blocked waiting for I/O completion.
 - ◆ The scheduler should schedule another job on the CPU.
- ▣ When the I/O completes
 - ◆ An interrupt is raised.
 - ◆ The OS moves the process from blocked back to the ready state.