

Operating Systems

Scheduling:
The Multi-Level Feedback Queue

Multi-Level Feedback Queue (MLFQ)

- ▣ A Scheduler that learns from the past to predict the future.
- ▣ Objective:
 - ◆ Optimize **turnaround time** → Run shorter jobs first
 - ◆ Minimize **response time** without *a priori knowledge of job length*.
- ◆ The multi-level feedback queue is an excellent example of a system that learns from the past to predict the future.

MLFQ: Basic Rules


- ▣ MLFQ has a number of distinct **queues**.
 - ◆ Each queue is assigned a different priority level.
- ▣ A job that is ready to run is on a single queue.
 - ◆ A job **on a higher queue** is chosen to run.
 - ◆ Use round-robin scheduling among jobs in the same queue

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

MLFQ: how the scheduler sets priorities?

- ▣ MLFQ varies the priority of a job based on **its observed behavior**.
- ▣ Example:
 - ◆ A job repeatedly relinquishes the CPU while waiting IOs → Keep its priority high
 - ◆ A job uses the CPU intensively for long periods of time → Reduce its priority.
- ▣ MLFQ will try to **learn** about processes as they run, and thus use the **history** of the job to *predict* its **future** behavior.

MLFQ Example



The scheduler would just alternate time slices between A and B because they are the highest priority jobs in the system

MLFQ: How to Change Priority

- ▣ we must keep in mind our workload:
 - ◆ a mix of interactive jobs that are short-running (and may frequently relinquish the CPU), and some longer-running “CPU-bound” jobs that need a lot of CPU time but where response time isn’t important

MLFQ: How to Change Priority


- MLFQ priority adjustment algorithm:

- ◆ **Rule 3:** When a job enters the system, it is placed at the highest priority
- ◆ **Rule 4a:** If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).
- ◆ **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level

In this manner, MLFQ approximates SJF

Example 1: A Single Long-Running Job

- A three-queue scheduler with time slice 10ms




Example 2: Along Came a Short Job

- Assumption:

- Job A: A long-running CPU-intensive job
- Job B: A short-running interactive job (20ms runtime)
- A has been running for some time, and then B arrives at time T=100.
- In this manner, MLFQ **approximates SJF**


because it doesn't **know** whether a job will be a short job or a long-running job, it first **assumes** it might be a short job, thus giving the job high priority. If it actually is a short job, it will run quickly and complete; if it is not a short job, it will slowly move down the queues, and thus soon prove itself to be a long-running process



Example 3: What About I/O?

- Assumption:

- Job A: A long-running CPU-intensive job
- Job B: An interactive job that need the CPU only for 1ms before performing an I/O





The MLFQ approach keeps an interactive job at the highest priority

Problems with the Basic MLFQ

- Starvation
 - ◆ If there are "too many" interactive jobs in the system.
 - ◆ Long-running jobs will never receive any CPU time.
- **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue.
 - ◆ First, processes are guaranteed not to starve
 - ◆ Second, if a CPU-bound job has become interactive, the scheduler treats it properly once it has received the priority boost.

The Priority Boost

- ◆ Example:
 - A long-running job(A) with two short-running interactive job(B, C)



Without(Left) and With(Right) Priority Boost

A: B: C: