

Operating Systems


Multiprocessor Scheduling

Multiprocessor Scheduling

- ❑ The rise of the **multicore processor** is the source of multiprocessor-scheduling proliferation.
 - ◆ **Multicore:** Multiple CPU cores are packed onto a single chip.
- ❑ Adding more CPUs does not make that single application run faster.
→ You'll have to rewrite application to run in parallel, using **threads**.

How to schedule jobs on **Multiple CPUs?**

Single CPU with cache




By keeping data in cache, the system can make slow memory appear to be a fast one


Cache coherence

- Consistency of shared resource data stored in multiple caches.

0. Two CPUs with caches sharing memory




1. CPU0 reads a data at address 1.




Cache coherence (Cont.)

2. D is updated and CPU1 is scheduled.



3. CPU1 re-reads the value at address A



CPU1 gets the **old value D** instead of the correct value D' .


Cache Affinity

- ▣ Keep a process on **the same CPU** if at all possible
 - ◆ A process builds up a fair bit of state in the cache of a CPU.
 - ◆ The next time the process runs, it will run faster if some of its state is *already present* in the cache on that CPU.

A multiprocessor scheduler should consider **cache affinity** when making its scheduling decision.

Single queue Multiprocessor Scheduling (SQMS)


- ▣ Put all jobs that need to be scheduled into **a single queue**.
 - ◆ Each CPU simply picks the next job from the globally shared queue.
 - ◆ Cons:
 - Some form of **locking** have to be inserted → **Lack of scalability**
 - **Cache affinity**
 - Example:



- Possible job scheduler across CPUs:

CPU0	A	E	D	C	B	... (repeat) ...
CPU1	B	A	E	D	C	... (repeat) ...
CPU2	C	B	A	E	D	... (repeat) ...
CPU3	D	C	B	A	E	... (repeat) ...

Scheduling Example with Cache affinity




- ◆ Preserving affinity for most
 - Jobs A through D are not moved across processors.
 - Only job E Migrating from CPU to CPU.
- ◆ Implementing such a scheme can be **complex**.

Multi-queue Multiprocessor Scheduling (MQMS)

- ▣ MQMS consists of **multiple scheduling queues**.
 - ◆ Each queue will follow a particular scheduling discipline.
 - ◆ When a job enters the system, it is placed on **exactly one** scheduling queue.
 - ◆ OS has to decide into which queue to place each job
 - ◆ Avoid the problems of information sharing and synchronization.

MQMS Example


- With **round robin**, the system might produce a schedule that looks like this:




MQMS provides more **scalability** and **cache affinity**.

Load Imbalance issue of MQMS

- After job C in Q0 finishes:




- After job A in Q0 finishes:



How to deal with load imbalance?

- The answer is to move jobs (**Migration**).

- ◆ Example:




How to deal with load imbalance? (Cont.)

- A more tricky case:



- A possible migration pattern:

- ◆ Keep switching jobs



Work Stealing

- ▣ Move jobs between queues
 - ◆ Implementation:
 - A source queue that is low on jobs is picked.
 - The source queue occasionally peeks at another target queue.
 - If the target queue is more full than the source queue, the source will “**steal**” one or more jobs from the target queue.
 - ◆ Cons:
 - *High overhead* and trouble *scaling*