

# **Operating Systems**

---

**Synchronization**

# Concurrency Example

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "common.h"
5  volatile int counter = 0;
6  int loops;
8  void *worker(void *arg) {
9      int i;
10     for (i = 0; i < loops; i++) {
11         counter++;
12     }
13     return NULL;
14 }
16 int
17 main(int argc, char *argv[])
18 {
19     loops = atoi(argv[1]);
24     pthread_t p1, p2;
25     printf("Initial value : %d\n", counter);
27     Pthread_create(&p1, NULL, worker, NULL);
28     Pthread_create(&p2, NULL, worker, NULL);
29     Pthread_join(p1, NULL);
30     Pthread_join(p2, NULL);
31     printf("Final value : %d\n", counter);
32     return 0;
33 }
```

# Concurrency Example (Cont.)

- loops determines how many times each of the two workers will **increment the shared counter** in a loop.
  - ◆ loops: 1000.

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- ◆ loops: 100000.

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

# Why is this happening?

- ▣ Increment a shared counter → take three instructions.
  1. Load the value of the counter from memory into register.
  2. Increment it
  3. Store it back into memory

counter = counter + 1

```
105      mov 0x8049a1c, %eax
108      add $0x1, %eax
113      mov %eax, 0x8049a1c
```

# Race condition

## Example with two threads

- counter = counter + 1 (default is 50)
- We expect the result is 52. However,

```
105    mov 0x8049a1c, %eax
108    add $0x1, %eax
113    mov %eax, 0x8049a1c
```

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
		<i>before critical section</i>			
		mov 0x8049a1c, %eax	100	0	50
		add \$0x1, %eax	105	50	50
			108	51	50
<b>interrupt</b>					
	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
<b>interrupt</b>					
	save T2's state				
	restore T1's state		108	51	50
		mov %eax, 0x8049a1c	113	51	<b>51</b>

# A few terminologies

- Race condition:
  - ◆ the results depend on the timing execution of the code.
  - ◆ Result is indeterminate.
- Critical section
  - ◆ A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread.
  - ◆ Multiple threads executing critical section can result in a race condition.
  - ◆ Need to support **atomicity** for critical sections (**mutual exclusion**)
  - ◆ This property guarantees that if one thread is executing within the critical section, the others will be prevented from doing so.

# The wish for atomicity

- ☐ Ideal approach; make the increment as a single assembly instruction

```
memory-add 0x8049alc, $0x1
```

- ☐ Atomically, in this context, means "as a unit", which sometimes we take as "all or none."
- ☐ In general, we do not have such instruction. Instead, we use lock.
- ☐ Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**).

```
1  lock_t mutex;
2  . . .
3  lock(&mutex);
4  balance = balance + 1; → Critical section
5  unlock(&mutex);
```

# Locks

---

# Locks: The Basic Idea

- ❑ Ensure that any **critical section** executes as if it were a single atomic instruction.

- ◆ An example: the canonical update of a shared variable

```
balance = balance + 1;
```

- ◆ Add some code around the critical section

```
1  lock_t mutex; // some globally-allocated lock 'mutex'  
2  ...  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

# Locks: The Basic Idea

- ▣ Lock variable holds the state of the lock.
  - ◆ **available** (or **unlocked** or **free**)
    - No thread holds the lock.
  - ◆ **acquired** (or **locked** or **held**)
    - Exactly one thread holds the lock and presumably is in a critical section.

# The semantics of the lock()

- ▣ lock()
  - ◆ Try to acquire the lock.
  - ◆ If no other thread holds the lock, the thread will **acquire** the lock.
  - ◆ **Enter** the *critical section*.
    - This thread is said to be the owner of the lock.
  - ◆ Other threads are *prevented from* entering the critical section while the first thread that holds the lock is in there.

# Pthread Locks - mutex

- ❑ The name that the POSIX library uses for a lock.

- ◆ Used to provide **mutual exclusion** between threads.

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()
4  balance = balance + 1;
5  Pthread_mutex_unlock(&lock);
```

- ◆ We may be using *different locks* to protect *different variables* → Increase **concurrency** (a more **fine-grained** approach).

# Building A Lock

- ▣ Efficient locks provided mutual exclusion at **low cost**.
- ▣ Building a lock need some help from the **hardware** and the **OS**.

# Evaluating locks – Basic criteria

## □ Mutual exclusion

- ◆ Does the lock work, preventing multiple threads from entering *a critical section*?

## □ Fairness

- ◆ Does each thread contending for the lock get a fair shot at acquiring it once it is free? (Starvation)

## □ Performance

- ◆ The time overheads added by using the lock
  - No contention
  - Contention of multiple threads on a single-core CPU
  - Contention multiple threads on multiple CPUs

# Controlling Interrupts

## ▫ **Disable Interrupts** for critical sections

- ◆ One of the earliest solutions used to provide mutual exclusion
- ◆ Invented for single-processor systems.

```
1 void lock() {  
2     DisableInterrupts();  
3 }  
4 void unlock() {  
5     EnableInterrupts();  
6 }
```

- ◆ Problem:
  - Require too much *trust* in applications
    - Greedy (or malicious) program could monopolize the processor.
  - Do not work on **multiprocessors**.
  - Code that masks or unmasks interrupts be executed *slowly* by modern CPUs.

# Why hardware support needed?

- ▣ **First attempt:** Using a *flag* denoting whether the lock is held or not.
  - ◆ The code below has problems.

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 → lock is available, 1 → held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11      mutex->flag = 1; // now SET it !
12  }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

# Why hardware support needed? (Cont.)

- ◆ **Problem 1:** No Mutual Exclusion (assume `flag=0` to begin)

Thread1

```
call lock()
while (flag == 1)
interrupt: switch to Thread 2
```

`flag = 1; // set flag to 1 (too!)`

Thread2

```
call lock()
while (flag == 1)
flag = 1;
interrupt: switch to Thread 1
```

- ◆ **Problem 2:** Spin-waiting wastes time waiting for another thread.
- ▣ So, we need an atomic instruction supported by **Hardware!**
  - ◆ *test-and-set* instruction, also known as *atomic exchange*

# Test And Set (Atomic Exchange)

- An atomic instruction to support the creation of simple locks

```
1 int TestAndSet(int *ptr, int new) {  
2     int old = *ptr;    // fetch old value at ptr  
3     *ptr = new;       // store 'new' into ptr  
4     return old;        // return the old value  
5 }
```

- ◆ **return**(testing) old value pointed to by the `ptr`.
- ◆ *Simultaneously update*(setting) said value to `new`.
- ◆ This sequence of operations is **performed atomically**.

# A Simple Spin Lock using test-and-set

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available,
7      // 1 that it is held
8      lock->flag = 0;
9  }
10
11 void lock(lock_t *lock) {
12     while (TestAndSet(&lock->flag, 1) == 1)
13         ;           // spin-wait
14 }
15
16 void unlock(lock_t *lock) {
17     lock->flag = 0;
18 }
```

```
1  int TestAndSet(int *ptr,
2                  int new) {
3      int old = *ptr;
4      *ptr = new;
5      return old;
6  }
```

- ◆ **Note:** To work correctly on *a single processor*, it requires a preemptive scheduler.

# Evaluating Spin Locks

- **Correctness:** yes

- ◆ The spin lock only allows a single thread to entry the critical section.

- **Fairness:** no

- ◆ Spin locks don't provide any fairness guarantees.
  - ◆ Indeed, a thread spinning may spin *forever*.

- **Performance:**

- ◆ In the single CPU, performance overheads can be quite *painful*.
  - ◆ If the number of threads roughly equals the number of CPUs, spin locks work *reasonably well*.

# Fetch-And-Add

- Atomically increment a value while returning the old value at a particular address.

```
1 int FetchAndAdd(int *ptr) {  
2     int old = *ptr;  
3     *ptr = old + 1;  
4     return old;  
5 }
```

**Fetch-And-Add Hardware atomic instruction (C-style)**

# Ticket Lock

- ▣ **Ticket lock** can be built with fetch-and add.

- ◆ Ensure progress for all threads. → **fairness**

```
1  typedef struct __lock_t {  
2      int ticket;  
3      int turn;  
4  } lock_t;  
5  
6  void lock_init(lock_t *lock) {  
7      lock->ticket = 0;  
8      lock->turn = 0;  
9  }  
10  
11 void lock(lock_t *lock) {  
12     int myturn = FetchAndAdd(&lock->ticket);  
13     while (lock->turn != myturn)  
14         ; // spin  
15 }  
16 void unlock(lock_t *lock) {  
17     FetchAndAdd(&lock->turn);  
18 }
```

```
1  int FetchAndAdd(int *ptr)  
2  {  
3      int old = *ptr;  
4      *ptr = old + 1;  
5      return old;  
6  }
```

# So Much Spinning

- ❑ Hardware-based spin locks are **simple** and they work.
- ❑ In some cases, these solutions can be quite **inefficient**.
  - ◆ Any time a thread gets caught *spinning*, it **wastes an entire time slice** doing nothing but checking a value.

**How To Avoid *Spinning*?**  
**We'll need OS Support too!**

# Using Queues: Sleeping Instead of Spinning

- ▣ **Queue** to keep track of which threads are waiting to enter the lock.
- ▣ park()
  - ◆ Put a calling thread to sleep
- ▣ unpark(threadID)
  - ◆ Wake a particular thread as designated by threadID.

# Using Queues: Sleeping Instead of Spinning

```
typedef struct __lock_t {
    int flag;    // lock is acquired or not
    int guard;   // to protect the queue
    queue_t *q;
} lock_t;
```

# Using Queues: Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

**Lock With Queues, Test-and-set, Yield, And Wakeup**

# Using Queues: Sleeping Instead of Spinning

```
22 void unlock(lock_t *m) {
23     while (TestAndSet(&m->guard, 1) == 1)
24         ; // acquire guard lock by spinning
25     if (queue_empty(m->q))
26         m->flag = 0; // let go of lock; no one wants it
27     else
28         unpark(queue_remove(m->q)); // hold lock (for next thread!)
29     m->guard = 0;
30 }
```

**Lock With Queues, Test-and-set, Yield, And Wakeup (Cont.)**

# Locks APIs

- Provide mutual exclusion to a critical section

- ◆ Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- ◆ Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and enter the critical section.
    - If another thread hold the lock → the thread will not return from the call until it has acquired the lock.

# Locks APIs (Cont.)

- All locks must be properly initialized.

- ◆ One way: using PTHREAD\_MUTEX\_INITIALIZER

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- ◆ The dynamic way: using pthread\_mutex\_init()

```
int rc = pthread_mutex_init(&lock, NULL);
assert(rc == 0); // always check success!
```

# Locks APIs (Cont.)

- Check errors code when calling lock and unlock

- ◆ An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```

- These two calls are used in lock acquisition

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_timelock(pthread_mutex_t *mutex,
                           struct timespec *abs_timeout);
```

- ◆ trylock: return failure if the lock is already held
  - ◆ timelock: return after a timeout

## Locks APIs (Cont.)

- These two calls are also used in **lock acquisition**

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timelock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

- trylock: return failure if the lock is already held
- timelock: return after a timeout or after acquiring the lock