

# **Operating Systems**

---

## **Condition Variables**

# Condition Variables

- ▣ There are many cases where a thread wishes to check whether a **condition** is true before continuing its execution.
- ▣ Example:
  - ◆ A parent thread might wish to check whether a child thread has *completed*.
  - ◆ This is often called a `join()`.

# Condition Variables

## A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3          // XXX how to indicate we are done?
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11         // XXX how to wait for child?
12         printf("parent: end\n");
13         return 0;
14     }
```

**What we would like to see here is:**

```
parent: begin
child
parent: end
```

# Parent waiting fore child: Spin-based Approach

```
1     volatile int done = 0;
2
3     void *child(void *arg) {
4         printf("child\n");
5         done = 1;
6         return NULL;
7     }
8
9     int main(int argc, char *argv[]) {
10        printf("parent: begin\n");
11        pthread_t c;
12        Pthread_create(&c, NULL, child, NULL); // create child
13        while (done == 0)
14            ; // spin
15        printf("parent: end\n");
16        return 0;
17    }
```

- ◆ This is hugely inefficient as the parent spins and **wastes CPU time**.

# How to wait for a condition

- ▣ Condition variable
  - ◆ Queue of threads
  - ◆ **Waiting** on the condition
    - An explicit queue that threads can put themselves on when some state of execution is not as desired.
  - ◆ **Signaling** on the condition
    - Some other thread, *when it changes said state*, can wake one of those waiting threads and allow them to continue.
- ▣ Three in a package
  - ◆ condition variable c
  - ◆ state variable done
  - ◆ lock m; // to protect state variable

# Definition and Routines

## ❑ Declare condition variable

```
Pthread_cond_t c;
```

- ◆ Proper initialization is required.

## ❑ Operation (the POSIX calls)

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);      // wait()
pthread_cond_signal(pthread_cond_t *c);                          // signal()
```

- ◆ The `wait()` call takes a mutex as a parameter.
  - The `wait()` call release the lock and put the calling thread to sleep.
  - When the thread wakes up, it must re-acquire the lock.

# Parent waiting for Child: Use a condition variable

```
1      int done = 0;
2      pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3      pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5      void thr_exit() {
6          Pthread_mutex_lock(&m);
7          done = 1;
8          Pthread_cond_signal(&c);
9          Pthread_mutex_unlock(&m);
10     }
11
12     void *child(void *arg) {
13         printf("child\n");
14         thr_exit();
15         return NULL;
16     }
17
18     void thr_join() {
19         Pthread_mutex_lock(&m);
20         while (done == 0)
21             Pthread_cond_wait(&c, &m);
22         Pthread_mutex_unlock(&m);
23     }
24
```

## Parent waiting for Child: Use a condition variable

(cont.)

```
25     int main(int argc, char *argv[]) {
26             printf("parent: begin\n");
27             pthread_t p;
28             Pthread_create(&p, NULL, child, NULL);
29             thr_join();
30             printf("parent: end\n");
31             return 0;
32 }
```

# Parent waiting for Child: Use a condition variable

## □ Parent:

- ◆ Creates the child thread and continues running itself.
- ◆ Calls into `thr_join()` to wait for the child thread to complete.
  - Acquires the lock.
  - Checks if the child is done.
  - Puts itself to sleep by calling `wait()`.
  - Releases the lock.

## □ Child:

- ◆ Prints the message "child".
- ◆ Calls `thr_exit()` to wake up the parent thread.
  - Grabs the lock.
  - Sets the state variable `done`.
  - Signals the parent thus waking it.

# A poor implementation

```
1     void thr_exit() {
2             done = 1;
3             Pthread_cond_signal(&c);
4     }
5
6     void thr_join() {
7         if (done == 0)
8             Pthread_cond_wait(&c);
9 }
```

- ◆ The issue here is a subtle **race condition**.
  - The parent calls `thr_join()`.
    - The parent checks the value of `done`.
    - It will see that it is 0 and try to go to sleep.
    - *Just before* it calls `wait` to go to sleep, the parent is interrupted and the child runs.
  - The child changes the state variable `done` to 1 and signals.
    - But no thread is waiting and thus no thread is woken.
    - When the parent runs again, it sleeps forever.

# The Producer / Consumer (Bound Buffer) Problem

## □ Producer

- ◆ Produce data items
- ◆ Wish to place data items in a buffer

## □ Consumer

- ◆ Grab data items out of the buffer consume them in some way

## □ Example: Multi-threaded web server

- ◆ *A producer* puts HTTP requests in to a work queue
- ◆ *Consumer threads* take requests out of this queue and process them

## Bounded buffer

- ▣ A bounded buffer is used when you pipe the output of one program into another.
  - ◆ Example: `grep foo file.txt | wc -l`
    - The `grep` process is the producer.
    - The `wc` process is the consumer.
    - Between them is an in-kernel bounded buffer.
  - ◆ Bounded buffer is Shared resource → **Synchronized access** is required.

# The Put and Get Routines (Version 1)

```
1      int buffer;
2      int count = 0;    // initially, empty
3
4      void put(int value) {
5          assert(count == 0);
6          count = 1;
7          buffer = value;
8      }
9
10     int get() {
11         assert(count == 1);
12         count = 0;
13         return buffer;
14     }
```

- ◆ Only put data into the buffer when `count` is zero.
  - i.e., when the buffer is *empty*.
- ◆ Only get data from the buffer when `count` is one.
  - i.e., when the buffer is *full*.

# Producer/Consumer Threads (Version 1)

```
1      void *producer(void *arg) {
2          int i;
3          int loops = (int) arg;
4          for (i = 0; i < loops; i++) {
5              put(i);
6          }
7      }
8
9      void *consumer(void *arg) {
10         int i;
11         while (1) {
12             int tmp = get();
13             printf("%d\n", tmp);
14         }
15     }
```

- ◆ **Producer** puts an integer into the shared buffer `loops` number of times.
- ◆ **Consumer** gets the data out of that shared buffer.

# Producer/Consumer: Single CV and If Statement

- ❑ Obviously the put() and get() routines have critical sections within them, as put() updates the buffer, and get() reads from it.
- ❑ A single condition variable `cond` and associated lock `mutex`

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              if (count == 1)                      // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                         // p4
11             Pthread_cond_signal(&cond);       // p5
12             Pthread_mutex_unlock(&mutex);     // p6
13         }
14     }
```

# Producer/Consumer: Single CV and If Statement

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex); // c1
20             if (count == 0) // c2
21                 Pthread_cond_wait(&cond, &mutex); // c3
22             int tmp = get(); // c4
23             Pthread_cond_signal(&cond); // c5
24             Pthread_mutex_unlock(&mutex); // c6
25             printf("%d\n", tmp);
26         }
27     }
```

- ◆ p1-p3: A producer waits for the buffer to be empty.
- ◆ c1-c3: A consumer waits for the buffer to be full.
- ◆ With just *a single producer* and *a single consumer*, the code works.

If we have **more than one of** producer and consumer?

# Thread Trace: Broken Solution (Version 1)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	$T_{c1}$ awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	$T_{c2}$ sneaks in ...
	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	$T_p$ awoken
	Ready	c6	Running		Ready	0	
c4	Running		Ready		Ready	0	Oh oh! No data

```

cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);           // p1
        if (count == 1)                      // p2
            Pthread_cond_wait(&cond, &mutex); // p3
        put(i);
        Pthread_cond_signal(&cond);          // p4
        Pthread_mutex_unlock(&mutex);         // p5
    }
}
  
```

```

16   void *consumer(void *arg) {
17       int i;
18       for (i = 0; i < loops; i++) {
19           Pthread_mutex_lock(&mutex); // c1
20           if (count == 0)             // c2
21               Pthread_cond_wait(&cond, &mutex); //
22           int tmp = get();          // c4
23           Pthread_cond_signal(&cond); // c5
24           Pthread_mutex_unlock(&mutex); // c6
25           printf("%d\n", tmp);
26       }
27   }
  
```

## Thread Trace: Broken Solution (Version 1)

- ▣ The problem arises for a simple reason:
  - ◆ After the producer woke  $T_{c1}$ , but before  $T_{c1}$  ever ran, the state of the bounded buffer *changed by  $T_{c2}$* .
  - ◆ There is no guarantee that when the woken thread runs, the state will still be as desired.

# Producer/Consumer: Single CV and While

- Consumer  $T_{c1}$  wakes up and **re-checks** the state of the shared variable.
  - If the buffer is empty, the consumer simply goes back to sleep.

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);                                // p1
8              while (count == 1)                                         // p2
9                  Pthread_cond_wait(&cond, &mutex);                      // p3
10             put(i);                                              // p4
11             Pthread_cond_signal(&cond);                            // p5
12             Pthread_mutex_unlock(&mutex);                           // p6
13         }
14     }
15 }
```

# Producer/Consumer: Single CV and While

(Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);                                // c1
20             while (count == 0)                                       // c2
21                 Pthread_cond_wait(&cond, &mutex);                  // c3
22             int tmp = get();                                       // c4
23             Pthread_cond_signal(&cond);                            // c5
24             Pthread_mutex_unlock(&mutex);                          // c6
25             printf("%d\n", tmp);
26         }
27     }
```

- ◆ A simple rule to remember with condition variables is to **always use while loops.**
- ◆ However, this code still has a bug (*next page*).

# Thread Trace: Broken Solution (Version 2)

	<b>State</b>		<b>State</b>		<b>State</b>	<b>Count</b>	<b>Comment</b>
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	grabs data
c5	Running		Ready		Sleep	0	<b>Oops! Woke</b>

## Thread Trace: Broken Solution (Version 2) (Cont.)

	State		State		State	Count	Comment
...	...	...	...	...	...	...	(cont.)
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	<b>Everyone asleep ...</b>

- ◆ A consumer should not wake **other** consumers, only producers, and vice-versa.

# The single Buffer Producer/Consumer Solution

- ❑ Use **two** condition variables and while

- ◆ **Producer** threads wait on the condition `empty`, and signals `fill`.
- ◆ **Consumer** threads wait on `fill` and signal `empty`.

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);
8              while (count == 1)
9                  Pthread_cond_wait(&empty, &mutex);
10             put(i);
11             Pthread_cond_signal(&fill);
12             Pthread_mutex_unlock(&mutex);
13         }
14     }
```

# The single Buffer Producer/Consumer Solution

(Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);
20             while (count == 0)
21                 Pthread_cond_wait(&fill, &mutex);
22             int tmp = get();
23             Pthread_cond_signal(&empty);
24             Pthread_mutex_unlock(&mutex);
25             printf("%d\n", tmp);
26         }
27     }
```