

# **Operating Systems**

## **Deadlock**

---

# Deadlock Bugs

Thread 1:

```
lock(L1);
```

```
lock(L2);
```

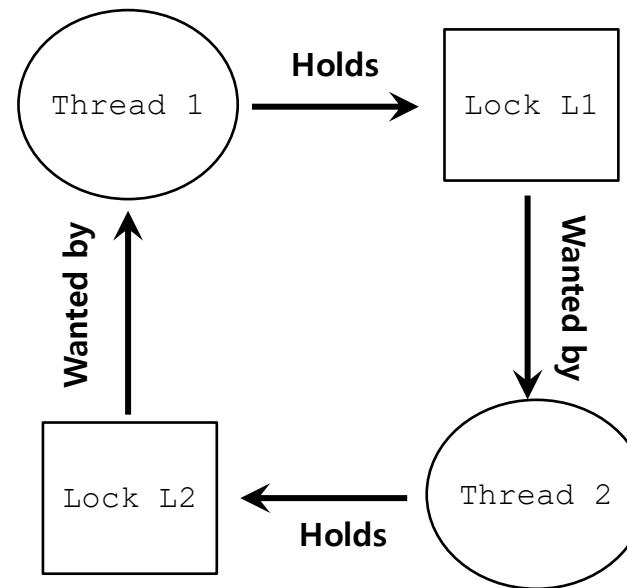
Thread 2:

```
lock(L2);
```

```
lock(L1);
```

- ◆ The presence of **a cycle**

- Thread1 is holding a lock L1 and waiting for another one, L2.
- Thread2 that holds lock L2 is waiting for L1 to be released.



# Why Do Deadlocks Occur?

- Reason 1:
  - ◆ In large code bases, **complex dependencies** arise between components.
- Reason 2:
  - ◆ Due to the nature of **encapsulation**
    - Hide details of implementations and make software easier to build in a modular way.
    - Such **modularity** *does not mesh well with locking.*

# Conditions for Deadlock

- Four conditions need to hold for a deadlock to occur.

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them.
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

- If any of these four conditions are not met, **deadlock cannot occur**.

# Solutions to Deadlock

- Prevention
- Avoidance
- Detect and Recover

# Prevention – Circular Wait

- Provide a total ordering on lock acquisition
  - ◆ This approach requires *careful design* of global locking strategies.
- Example:
  - ◆ There are two locks in the system (L1 and L2)
  - ◆ We can prevent deadlock by always acquiring L1 before L2.

# Prevention – Hold-and-wait

- Acquire all locks **at once, atomically.**

```
1  lock(prevention);  
2  lock(L1);  
3  lock(L2);  
4  ...  
5  unlock(prevention);
```

- This code guarantees that **no untimely thread switch can occur *in the midst of* lock acquisition.**
- Problem:**
  - Require us to know when calling a routine exactly which locks must be held and to acquire them ahead of time.
  - Decrease *concurrency* as all locks must be acquired early on (at once) instead of when they are truly needed.

# Prevention – No Preemption

- ▣ **Multiple lock acquisition** often gets us into trouble because when waiting for one lock **we are holding another**.
- ▣ `trylock()`
  - ◆ Used to build a *deadlock-free, ordering-robust* lock acquisition protocol.
  - ◆ Grab the lock (if it is available).
  - ◆ Or, return -1: you should try again later.

```
1  top:  
2      lock(L1);  
3      if( tryLock(L2) == -1 ) {  
4          unlock(L1);  
5          goto top;  
6      }
```

# Prevention – No Preemption (Cont.)

## ❑ livelock

- ◆ Both systems are running through the code sequence *over and over again*.
- ◆ Progress is not being made.
- ◆ Solution:
  - Add **a random delay** before looping back and trying the entire thing over again.

# Prevention – Mutual Exclusion

## ▫ wait-free

- ◆ Using powerful **hardware instruction**.
- ◆ You can build data structures in a manner that *does not require explicit locking*.

```
1 int CompareAndSwap(int *address, int expected, int new) {  
2     if (*address == expected) {  
3         *address = new;  
4         return 1; // success  
5     }  
6     return 0;  
7 }
```

# Prevention – Mutual Exclusion (Cont.)

## □ Example: list insertion

- ◆ Surrounding this code with a **lock acquire** and **release**.

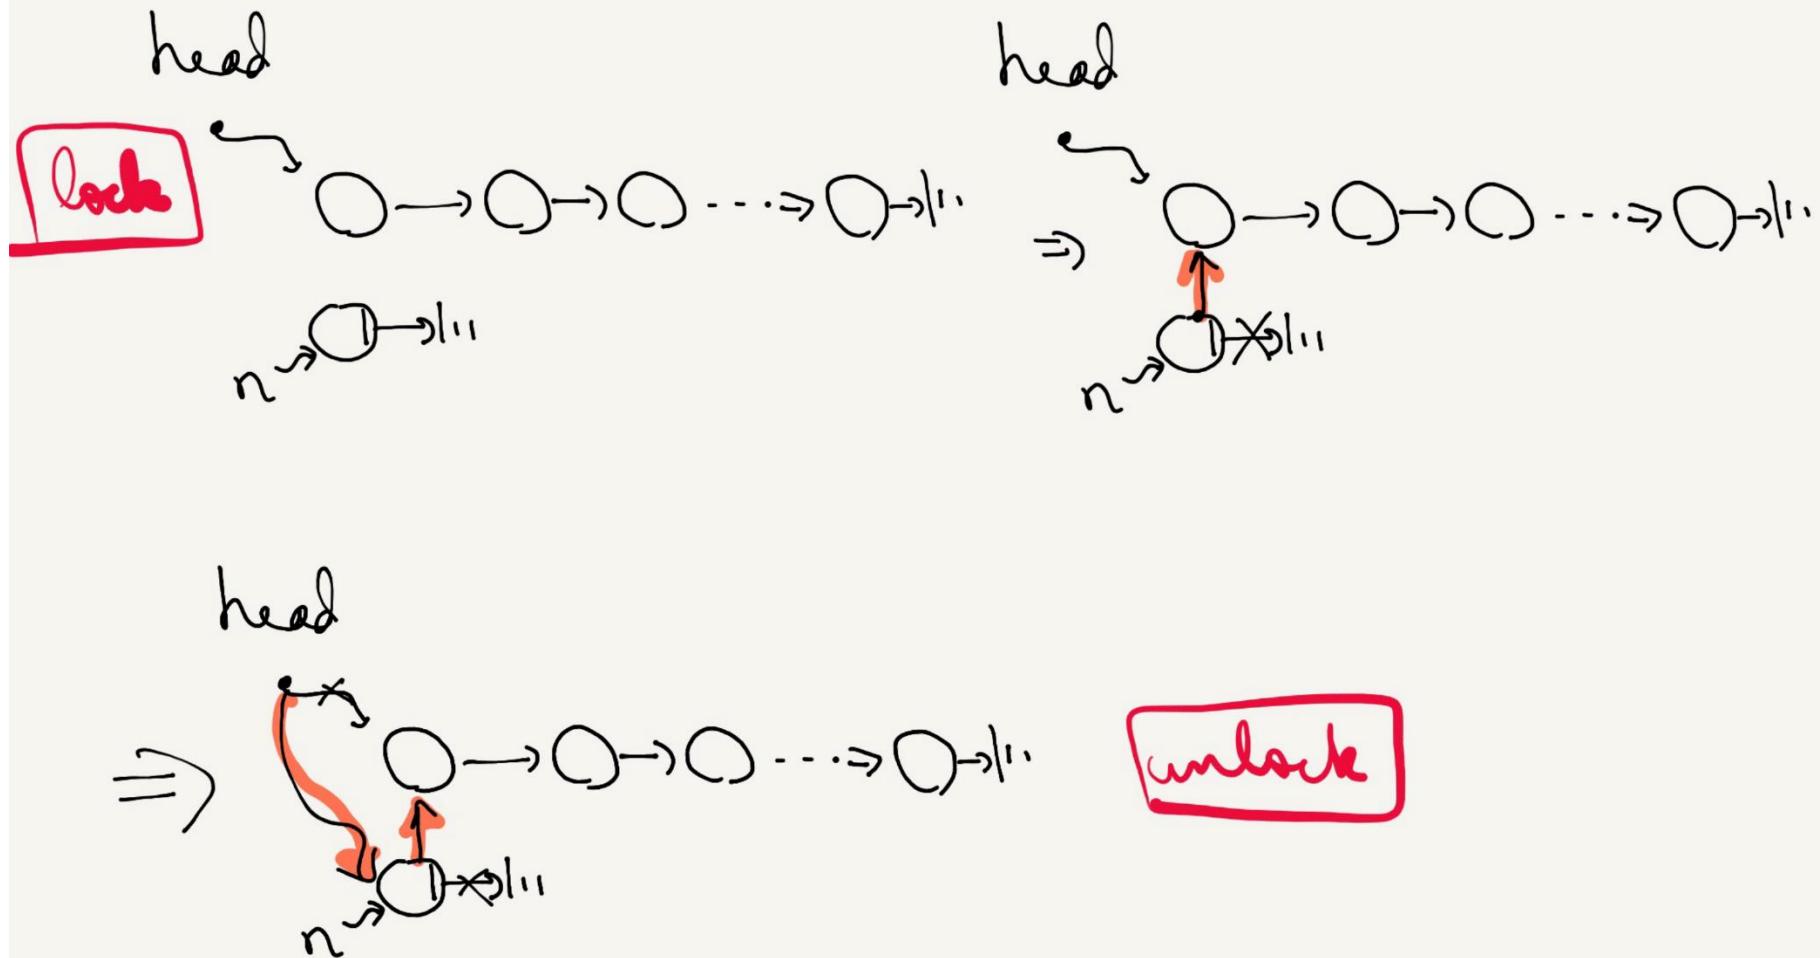
```
1 void insert(int value){  
2     node_t * n = malloc(sizeof(node_t));  
3     assert( n != NULL );  
4     n->value = value ;  
5     lock(listlock); // begin critical section  
6     n->next    = head;  
7     head       = n;  
8     unlock(listlock) ; //end critical section  
9 }
```

- ◆ **wait-free manner** using the compare-and-swap instruction

```
1 int CompareAndSwap(int *address, int expected, int new) {  
2     if(*address == expected) {  
3         *address = new;  
4         return 1; // success  
5     }  
6     return 0;  
7 }
```

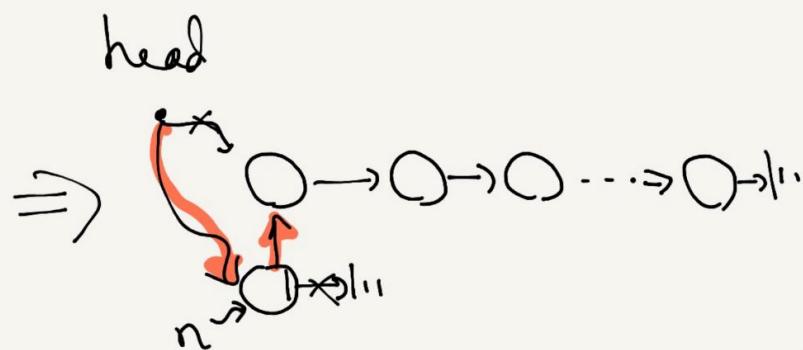
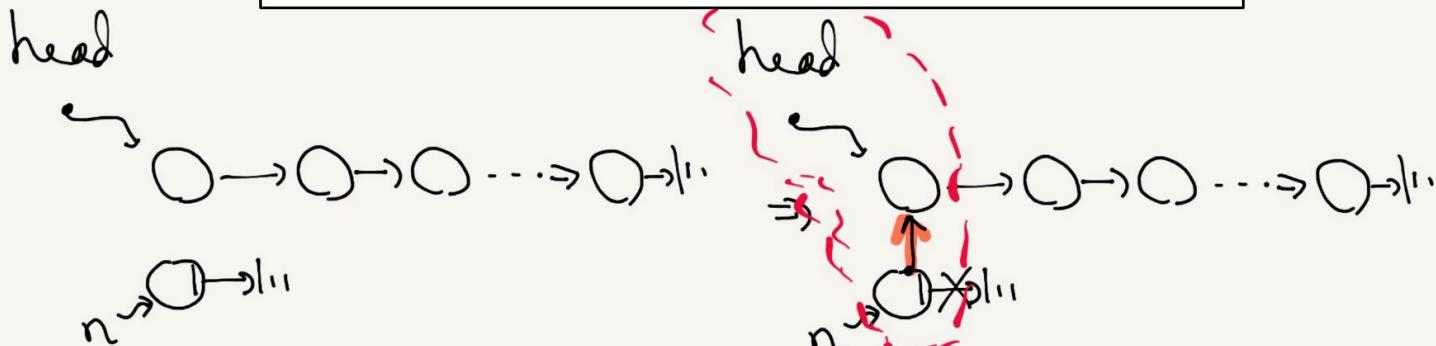
```
1 void insert(int value) {  
2     node_t *n = malloc(sizeof(node_t));  
3     assert(n != NULL);  
4     n->value = value;  
5     do {  
6         n->next = head;  
7     } while (CompareAndSwap(&head, n->next, n));  
8 }
```

## Insert with lock



# Lock-free insert

```
1 void insert(int value) {  
2     node_t *n = malloc(sizeof(node_t));  
3     assert(n != NULL);  
4     n->value = value;  
5     do {  
6         n->next = head;  
7     } while (CompareAndSwap(&head, n->next, n));  
8 }
```



Compare And Swap    if ( $\text{head} == n \rightarrow \text{next}$ )  
                             $\text{head} = n;$

# Deadlock Avoidance via Scheduling

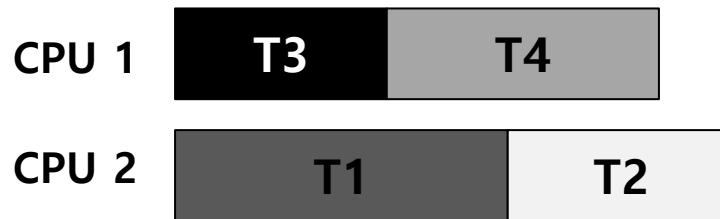
- ▣ Deadlock Avoidance
  - ◆ Get the information about the locks various threads might grab during their execution.
  - ◆ schedule the threads in a way to guarantee no deadlock can occur.
- ▣ In some scenarios, **deadlock avoidance** is preferable.
- ▣ Problem: Global knowledge is required.

# Example of Deadlock Avoidance via Scheduling (1)

- We have two processors and four threads.
  - ◆ Lock acquisition demands of the threads:

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no

- ◆ A smart scheduler could compute that as long as T1 and T2 are not run at the same time, no deadlock could ever arise.

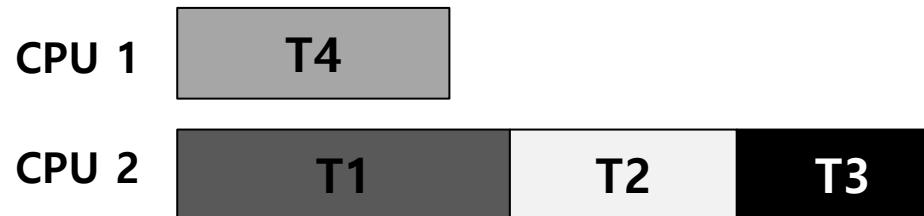


## Example of Deadlock Avoidance via Scheduling (2)

- More contention for the same resources

	T1	T2	T3	T4
L1	yes	yes	yes	no
L2	yes	yes	yes	no

- A possible schedule that guarantees that *no deadlock* could ever occur.



- The total time to complete the jobs is lengthened considerably.

# Detect and Recover

- ▣ Allow deadlock to occasionally occur and then *take some action.*
  - ◆ Example: if an OS froze, you would reboot it.
- ▣ View system as graph
  - ◆ Processes and Resources are nodes
  - ◆ Resource Requests and Assignments are edges
- ▣ If graph has no cycles → no deadlock
- ▣ If graph contains a cycle
  - ◆ Definitely deadlock if only one instance per resource
  - ◆ Otherwise, maybe deadlock, maybe no
- ▣ Many database systems employ *deadlock detection* and *recovery technique*.
  - ◆ A deadlock detector **runs periodically**.
  - ◆ Building a **resource graph** and checking it for cycles.
  - ◆ In deadlock, the system **need to be restarted**.