

# **Operating Systems**

---

## **File System**

# Persistent Storage


- ▣ Keep a data **intact** even if there is a power loss.
  - ◆ Hard disk drive
  - ◆ Solid-state storage device
- ▣ Two key abstractions in the virtualization of storage
  - ◆ File
  - ◆ Directory

- ▣ A linear array of bytes
- ▣ Each file has low-level name as **inode number**
  - ◆ The user is not aware of this name.
- ▣ A file system is the software that manages how data is stored, organized, and accessed on a storage device.
- ▣ Filesystem has a responsibility to store data persistently on disk.

# Directory

- ▣ Directory is like a file, also has a low-level name.
  - ◆ It contains a list of (user-readable name, low-level name) pairs.
  - ◆ Each entry in a directory refers to either *files* or other *directories*.
- ▣ Example)
  - ◆ A directory has an entry ("foo", "10")
    - A file "foo" with the low-level name "10"

# Directory Tree (Directory Hierarchy)



An Example Directory Tree

**Valid files (absolute pathname) :**

/foo/bar.txt  
/bar/foo/bar.txt

**Valid directory :**

/  
/foo  
/bar  
/bar/bar  
/bar/foo/









} Sub-directories

# The Way To Think

- ▣ There are two different aspects to implement file system
  - ◆ **Data structures**
    - What types of on-disk structures are utilized by the file system to organize its data and metadata?
  - ◆ **Access methods**
    - How does it map the calls made by a process as `open()`, `read()`, `write()`, etc.
    - Which structures are read during the execution of a particular system call?


# Overall Organization

- Let's develop the overall organization of the file system data structure.
- Divide the disk into **blocks**.
  - Block size is 4 KB.
  - The blocks are addressed from  $0$  to  $N - 1$ .



# Data region in file system

- Reserve **data region** to store user data




- File system has to track which data block comprise a file, the size of the file, its owner, etc.
- An inode is a data structure that stores all the metadata of a file except its name.

How we store these **inodes** in file system?


# Inode table in file system

- ❑ An inode table is a collection of inodes that store all the metadata about files on a disk.
- ❑ Reserve some space for **inode table**
  - ◆ This holds an array of on-disk inodes.
  - ◆ Ex) inode tables : 3 ~ 7, inode size : 256 bytes
    - 4-KB block can hold 16 inodes.
    - The filesystem contains 80 inodes. (maximum number of files)




# Allocation structures

- This is to track whether inodes or data blocks are free or allocated.
- Use **bitmap**, each bit indicates free(0) or in-use(1)
  - ◆ **data bitmap**: for data region
  - ◆ **inode bitmap**: for inode table



# Superblock

- Super block contains this **information** for particular file system
    - ◆ Ex) The number of inodes, begin location of inode table. etc



- ◆ Thus, when mounting a file system, OS will read the superblock first, to initialize various information.

# File Organization: The inode

- Each inode is referred to by inode number.
  - by inode number, File system calculate where the inode is on the disk.
  - Ex) inode number: 32
    - Calculate the offset into the inode region ( $32 \times \text{sizeof(inode)}$ ) (256 bytes) = 8192
    - Add start address of the inode table(12 KB) + inode region(8 KB) = 20 KB

The Inode table

				iblock 0				iblock 1				iblock 2				iblock 3				iblock 4			
Super	i-bmap	d-bmap	0	1	2	3	16	17	18	19	32	33	34	35	48	49	50	51	64	65	66	67	
			4	5	6	7	20	21	22	23	36	37	38	39	52	53	54	55	68	69	70	71	
			8	9	10	11	24	25	26	27	40	41	42	43	56	57	58	59	72	73	74	75	
			12	13	14	15	28	29	30	31	44	45	46	47	60	61	62	63	76	77	78	79	
0KB	4KB	8KB	12KB	16KB	20KB	24KB	28KB	32KB															

## File Organization: The inode (Cont.)

- ▣ inode have all of the information about a file
  - ◆ File type (regular file, directory, etc.),
  - ◆ Size, the number of blocks allocated to it.
  - ◆ Protection information(who owns the file, who can access, etc).
  - ◆ Time information.
  - ◆ Etc.

# File Organization: The inode (Cont.)

---

Size	Name	
2	mode	What is this inode field for? can this file be read/written/executed?
2	uid	who owns this file?
4	size	how many bytes are in this file?
4	time	what time was this file last accessed?
4	ctime	what time was this file created?
4	mtime	what time was this file last modified?
4	dtime	what time was this inode deleted?
4	gid	which group does this file belong to?
2	links_count	how many hard links are there to this file?
2	blocks	how many blocks have been allocated to this file?
4	flags	how should ext2 use this inode?
4	osd1	an OS-dependent field
<b>60</b>	<b>block</b>	<b>a set of disk pointers (15 total)</b>
4	generation	file version (used by NFS)
4	file_acl	a new permissions model beyond mode bits
4	dir_acl	called access control lists
4	faddr	an unsupported field
12	i_osd2	another OS-dependent field

## The EXT2 Inode

# The Multi-Level Index

- ▣ To support bigger files, we use multi-level index.
- ▣ **Indirect pointer** points to a block that contains more pointers.
  - ◆ inode have fixed number of direct pointers (12) and a single indirect pointer.
- ▣ The number of pointers per indirect block is:

$$\frac{4kB}{4} = 1024 \text{ pointers}$$

- ◆ If a file grows large enough, an indirect block is allocated, inode's slot for an indirect pointer is set to point to it.
  - $(12 + 1024) \times 4 \text{ K or } 4144 \text{ KB}$

## The Multi-Level Index (Cont.)

- ▣ Double indirect pointer points to a block that contains indirect blocks.
  - ◆ Allow file to grow with an additional  $1024 \times 1024$  or 1 million 4KB blocks.
- ▣ Multi-Level Index approach to pointing to file blocks.
  - ◆ Ex) twelve direct pointers, a single and a double indirect block.
    - over 4GB in size  $(12+1024+1024^2) \times 4KB$
- ▣ Many file system use a multi-level index.
  - ◆ Linux EXT2, EXT3, NetApp's WAFL, Unix file system.

# The Multi-Level Index (Cont.)

**Most files are small**

**Average file size is growing**

**Most bytes are stored in large files**

**File systems contains lots of files**

**File systems are roughly half full**

**Directories are typically small**

Roughly 2K is the most common size

Almost 200K is the average

A few big files use most of the space

Almost 100K on average

Even as disks grow, file system remain ~50% full

Many have few entries; most have 20 or fewer

## File System Measurement Summary

# Directory Organization

- Directory contains a list of (entry name, inode number) pairs.
- Each directory has two extra files **."dot"** for current directory and **.."dot-dot"** for parent directory
  - ◆ For example, `dir` has three files (`foo`, `bar`, `foobar`)

<b>inum</b>	<b>reclen</b>	<b>strlen</b>	<b>name</b>
5	4	2	.
2	4	3	..
12	4	4	foo
13	4	4	bar
24	8	7	foobar

**on-disk for dir**

# Free Space Management

- ▣ File system track which inode and data block are free or not.
- ▣ In order to manage free space, we have two simple bitmaps.
  - ◆ When file is newly created, it allocated inode by searching the inode bitmap and update on-disk bitmap.
  - ◆ Pre-allocation policy is commonly used for allocate contiguous blocks.
    - some Linux file systems, such as ext2 and ext3, will look for a sequence of blocks (say 8) that are free when a new file is created and needs data blocks

# Access Paths: Reading a File From Disk

- ▣ Issue an `open ("/foo/bar", O_RDONLY)`,
  - ◆ Traverse the pathname and thus locate the desired inode.
  - ◆ Begin at the root of the file system (/)
    - In most Unix file systems, the root inode number is 2
  - ◆ Filesystem reads in the block that contains inode number 2.
  - ◆ Look inside of it to find pointer to data blocks (contents of the root).
  - ◆ By reading in one or more directory data blocks, It will find "foo" directory.
  - ◆ Traverse recursively the path name until the desired inode ("bar")
  - ◆ Check final permissions, allocate a file descriptor for this process and returns file descriptor to user.

## Access Paths: Reading a File From Disk (Cont.)

- ▣ Issue `read()` to read from the file.
  - ◆ Read in the first block of the file, consulting the inode to find the location of such a block.
    - Update the inode with a new last accessed time.
    - Update in-memory open file table for file descriptor, the file offset.
- ▣ When file is closed:
  - ◆ File descriptor should be deallocated, but for now, that is all the file system really needs to do. No disk I/Os take place.

# Access Paths: Reading a File From Disk (Cont.)

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
<code>open(bar)</code>			read	read	read	read	read			
<code>read()</code>				read				read		
<code>read()</code>				read					read	
<code>read()</code>				read						read

**File Read Timeline (Time Increasing Downward)**

## Access Paths: Writing to Disk

- ▣ Issue `write()` to update the file with new contents.
- ▣ File may allocate a block (unless the block is being overwritten).
  - ◆ Need to update data block, data bitmap.
  - ◆ It generates five I/Os:
    - one to read the data bitmap
    - one to write the data bitmap (to reflect its new state to disk)
    - two more to read and then write the inode
    - one to write the actual block itself.
  - ◆ To create file, it also allocate space for directory, causing high I/O traffic.

# Access Paths: Writing to Disk (Cont.)

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
create (/foo/bar)		read write	read	read	read write	read	read	write		
write()	read write			read				write		
write()	read write			read					write	
write()	read write			read						write

File Creation Timeline (Time Increasing Downward)

# Caching and Buffering

- ▣ Reading and writing files are expensive, incurring many I/Os.
  - ◆ For example, long pathname(/1/2/3/..../100/file.txt)
    - One to read the inode of the directory and at least one read its data.
    - Literally perform hundreds of reads just to open the file.
- ▣ In order to reduce I/O traffic, file systems aggressively use system memory(DRAM) to cache.
- ▣ Read I/O can be avoided by large cache.